

COL106

Data Structures and Algorithms

Subodh Sharma and Rahul Garg

Fibonacci Heaps

Based on slides by: Kevin Wayne, Princeton University, Data Structures, Stanford University

Why Binomial/Fibonacci Heaps?

Operation	Linked List	Binary Heap	Binomial Heap	Fibonacci Heap [†]	Relaxed Heap
<i>make-heap</i>	1	1	1	1	1
<i>is-empty</i>	1	1	1	1	1
<i>insert</i>	1	$\log n$	$\log n$	1	1
<i>delete-min</i>	n	$\log n$	$\log n$	$\log n$	$\log n$
<i>decrease-key</i>	n	$\log n$	$\log n$	1	1
<i>delete</i>	n	$\log n$	$\log n$	$\log n$	$\log n$
<i>union</i>	1	n	$\log n$	1	1
<i>find-min</i>	n	1	$\log n$	1	1

n = number of elements in priority queue

Runtime of
Dijkstra's/Prim's
Algorithm

$O(|V|^2)$

$O(|E|\log(|V|))$

$O(|E| + |V|\log(|V|))$

Dijkstra's/Prim's

1 make-heap

$|V|$ insert

$|V|$ delete-min

$|E|$ decrease

Why Fibonacci Heaps?

- History. [\[Fredman and Tarjan, 1986\]](#)
 - Ingenious data structure and analysis.
 - Original motivation: improve Dijkstra's shortest path algorithm from $O(E \log V)$ to $O(E + V \log V)$
 - Also works for Prim's MST algorithm
- $O(1)$ decrease key operation (amortized)
- Need all your attention

Fibonacci Heaps: Key Ideas

- Binomial trees
- Lazy merging of trees at the heap root
 - Only merge trees during deleteMin operation
 - Union of two heaps in $O(1)$ time
- Amortized analysis
- Decrease key implemented using direct cutting of the node subtree and moving to root
 - No longer perfect binomial trees
 - Maintain nodes as marked and limit such imperfections
 - Nice properties (degrees, size) of binomial tree are preserved

Amortized Analysis

- Consider an algorithm taking times t_1, t_2, \dots, t_k for performing k operations in sequence
- We cannot bound the worst-case time t_i
- Maintain a potential function $\phi(i)$ at every step
- $\phi(i)$ is like a computation bank.
 - You may deposit and withdraw from it to bound worst case amortized time at_i
 - $\phi(i)$ may depend on the internal state of the data structures
- Amortized time at step i , $at_i = t_i + \phi(i+1) - \phi(i)$
- Total time = $\sum_{i=1}^T t_i = \sum_{i=1}^T at_i - \phi(T+1) + \phi(1)$
- If $\phi(T+1) - \phi(1)$, then total time is same as total amortized time
- Example: Analysis of dynamic arrays

Example: Time to Increment a Binary Number

- Given an n-bit binary number
- How long does it take to increment it?

$b_n b_{n-1} \dots b_2 b_1 b_0$

Worst case time: $O(n)$ because if $b_{n-1} \dots b_0$ are all 1's then one needs to flip all the bits

Let ϕ = number of bits that are 1

Amortized time:


- The algo start with flipping the lsb's that are 1
- Stops when it reaches a bit that is 0 after making it 1
- Time taken: Number of 1's flopped to zero + 1
- Amortized time taken: 2

Fibonacci Heaps: Lazy Merging at the Root

Fibonacci Heaps: Lazy Merging at the Root

Basic idea.

- Similar to binomial heaps, but less rigid structure.
- Binomial heap: **eagerly** consolidate trees after each *insert*.



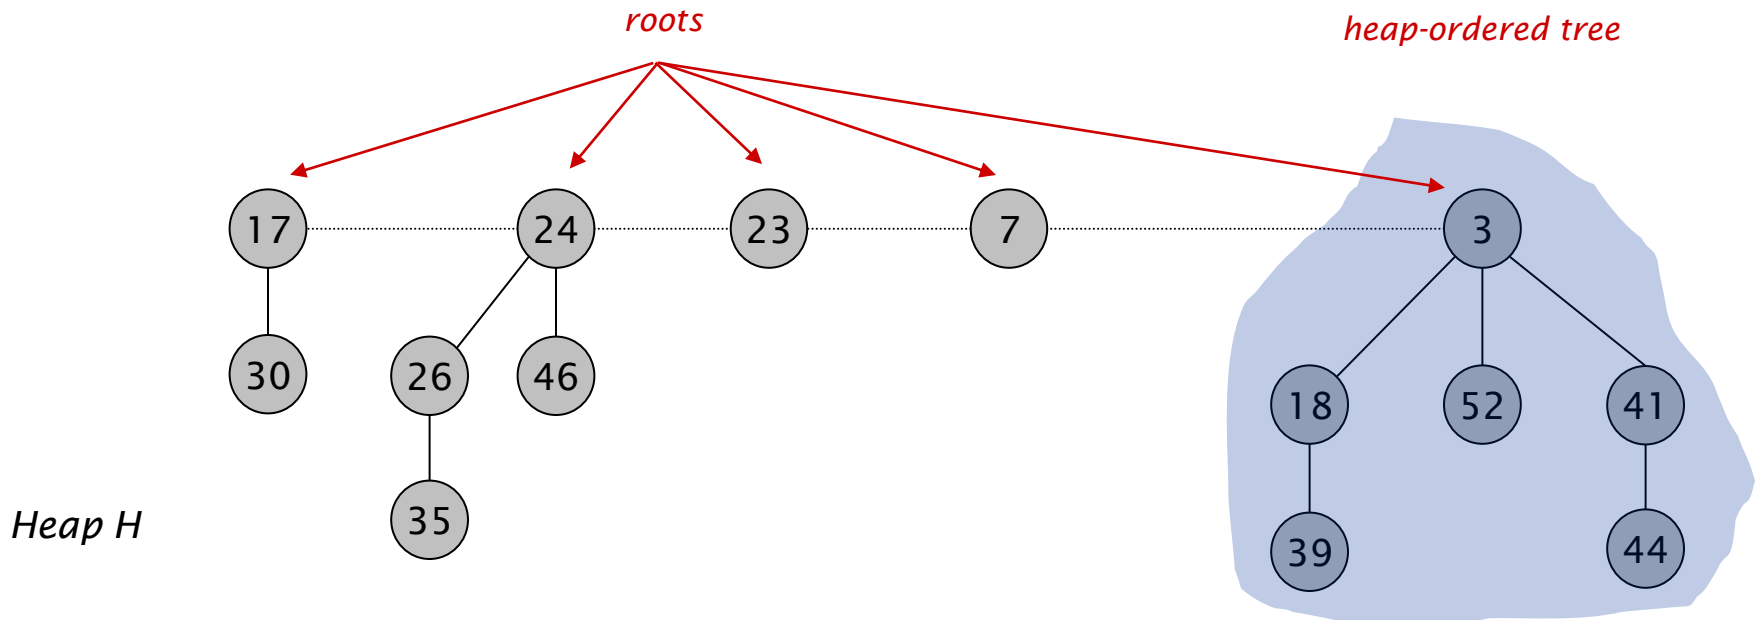
- Fibonacci heap: **lazily** defer consolidation until next *delete-min*.

Fibonacci Heaps: Structure

Fibonacci heap.

- Set of **heap-ordered** trees.
- Maintain pointer to minimum element.
- Set of marked nodes.

each parent larger than its children




Fibonacci Heaps: Structure

Fibonacci heap.

- Set of heap-ordered trees.
- **Maintain pointer to minimum element.**
- Set of marked nodes.

find-min takes $O(1)$ time




Fibonacci Heaps: Structure


Fibonacci heap.

- Set of heap-ordered trees.
- **Maintain pointer to minimum element.**
- Set of marked nodes.


merge heaps takes $O(1)$ time



Binomial Heap



Fibonacci Heap



~~Zero or one copies of binomial trees B_k~~

Can have many more trees

No longer binomial trees, but “defective” binomial trees


All trees satisfy the heap property

Fib Trees

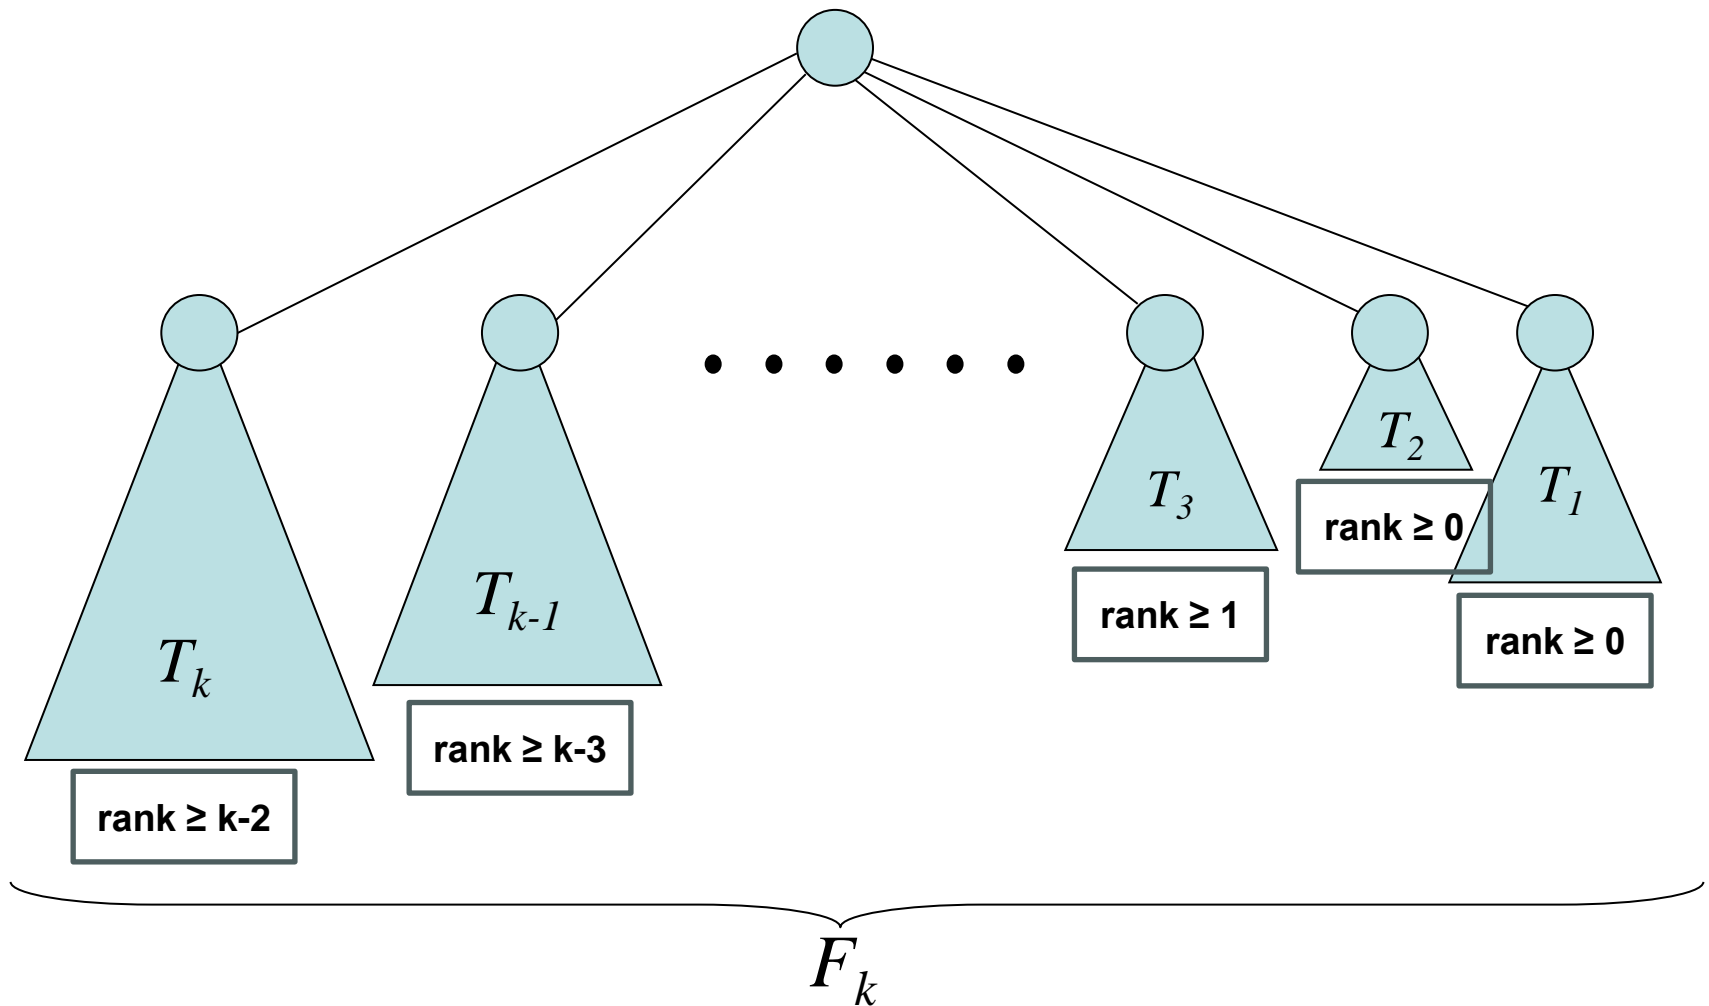
- Define rank of a tree, $\text{rank}(T)$ as the number of children of the root of the tree T
- The Fib F_k is a tree of rank(k) defined recursively
 - F_0 Consists of a single node
 - F_1 A single node with a single child
 - \vdots
 - F_k is a Fib tree of rank k with children T_1, T_2, \dots, T_k such that

$$\text{rank}(T_i) \geq \begin{cases} 0, & \text{if } i = 1 \\ i - 2, & \text{if } i \geq 2 \end{cases}$$

Binomial Trees



Fib Trees



Properties of Fib Trees

Define: $\text{size}(T)$ as the number of nodes in T

$D(T) = \max$ degree of the nodes in T

Lemma: For the Fib tree F_k

1. $\text{size}(F_k) \geq \phi^k$ where $\phi = (1 + \sqrt{5}) / 2$
2. $D(F_k) \leq \log_{\phi}(\text{size}(F_k))$


Why Fib Trees?

- **Goal:** Implement decrease-key in amortized time $O(1)$

Why Fib Trees?

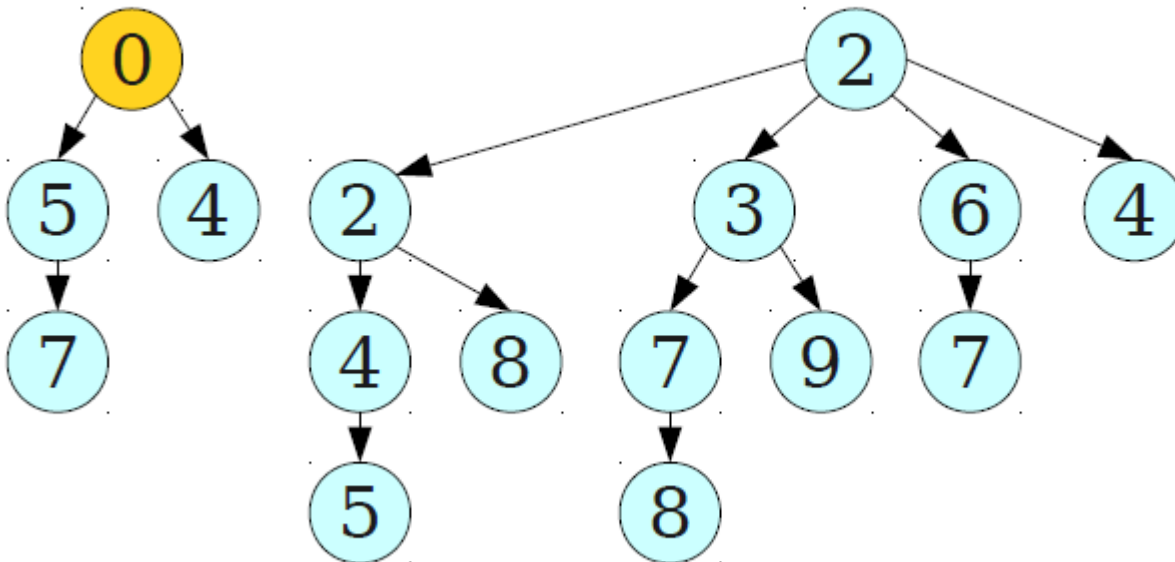
- **Goal:** Implement decreaseKey in amortized time $O(1)$

- Naive implementation
- Compare with the parent and bubble up
- May need $O(\log(n))$ time



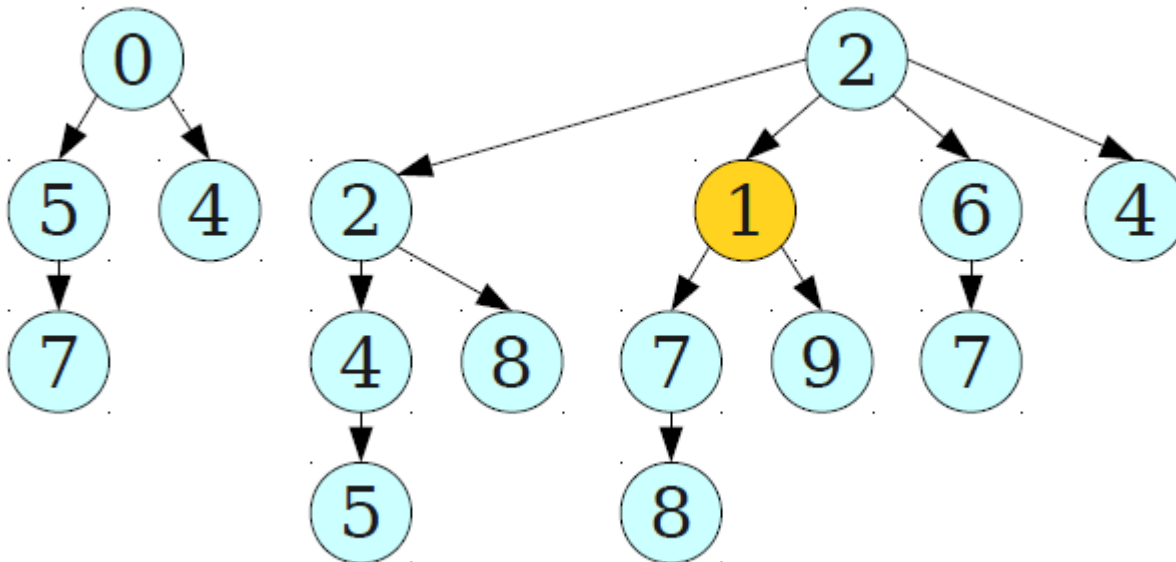
Our Idea

- **Goal:** Implement decreaseKey in amortized time $O(1)$



Our Idea

- **Goal:** Implement decreaseKey in amortized time $O(1)$

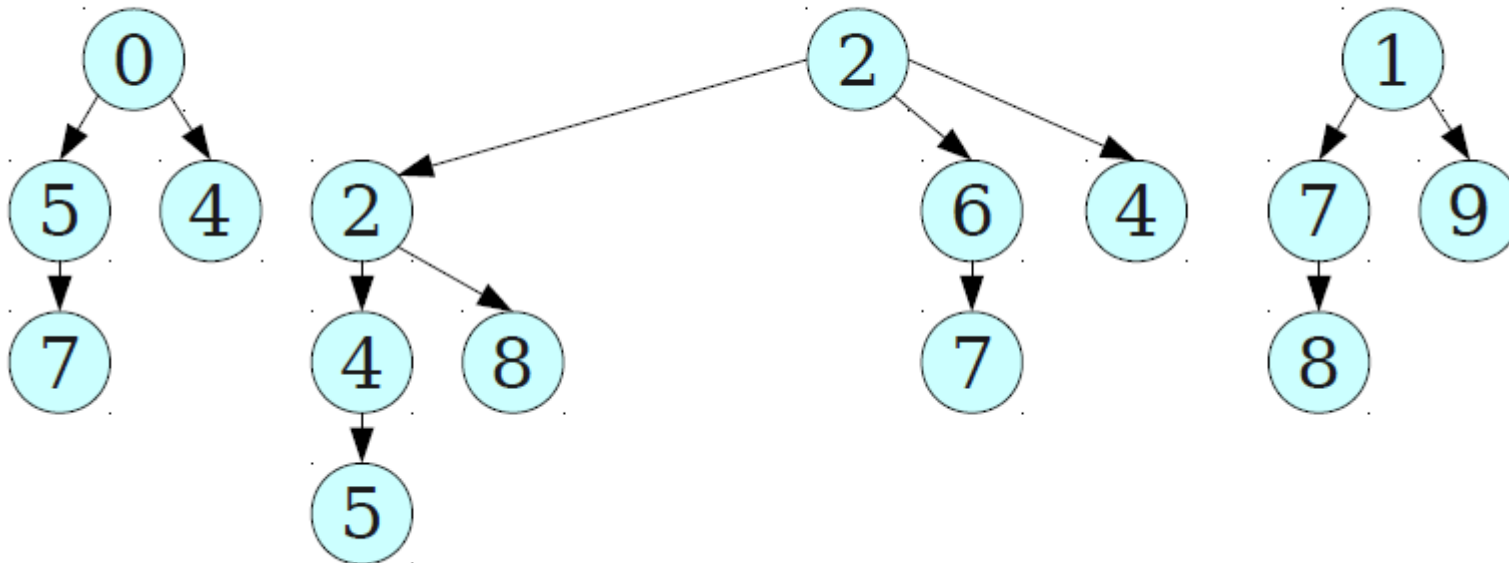


Our Idea

- **Goal:** Implement decreaseKey in amortized time $O(1)$

decreaseKey(x)

- Cut the node x and move subtree to root
- $O(1)$ operation



decreaseKey

- To implement *decrease-key* efficiently:
 - Lower the key of the specified node
 - If its key is greater than or equal to its parent's key, we're done
 - Otherwise, cut that node from its parent and hoist it up to the root list, optionally updating the min pointer
 - Time required: $O(1)$
- This requires some changes to the tree representation

Fibonacci Heaps: Structure


Fibonacci heap.

- Set of heap-ordered trees.
- Maintain pointer to minimum element.
- Set of marked nodes.

↖ If a node's child has been cut it is marked

No node can have more than one child cut


Used to keep heaps flat (stay tuned)



Fibonacci Heaps: Notation

Notation.

- n = number of nodes in heap.
- $rank(x)$ = number of children of node x .
- $rank(H)$ = max rank of any node in heap H .
- $trees(H)$ = number of trees in heap H .
- $marks(H)$ = number of marked nodes in heap H .



Fibonacci Heaps: Potential Function

$$\Phi(H) = \text{trees}(H) + 2 \cdot \text{marks}(H)$$

potential of heap H


trees(H) = 5

marks(H) = 3

$\Phi(H) = 5 + 2 \cdot 3 = 11$

min

Heap H




Insert

Fibonacci Heaps: Insert

Insert.

- Create a new singleton tree.
- Add to root list; update min pointer (if necessary).

insert 21




Fibonacci Heaps: Insert

Insert.

- Create a new singleton tree.
- Add to root list; update min pointer (if necessary).

insert 21



Fibonacci Heaps: Insert Analysis


Actual cost. $O(1)$

Change in potential. $+1$

Amortized cost. $O(1)$

$$\Phi(H) = \text{trees}(H) + 2 \cdot \text{marks}(H)$$

potential of heap H



Delete Min

Delete Min

We merge all the trees in the Heap in the deleteMin operation

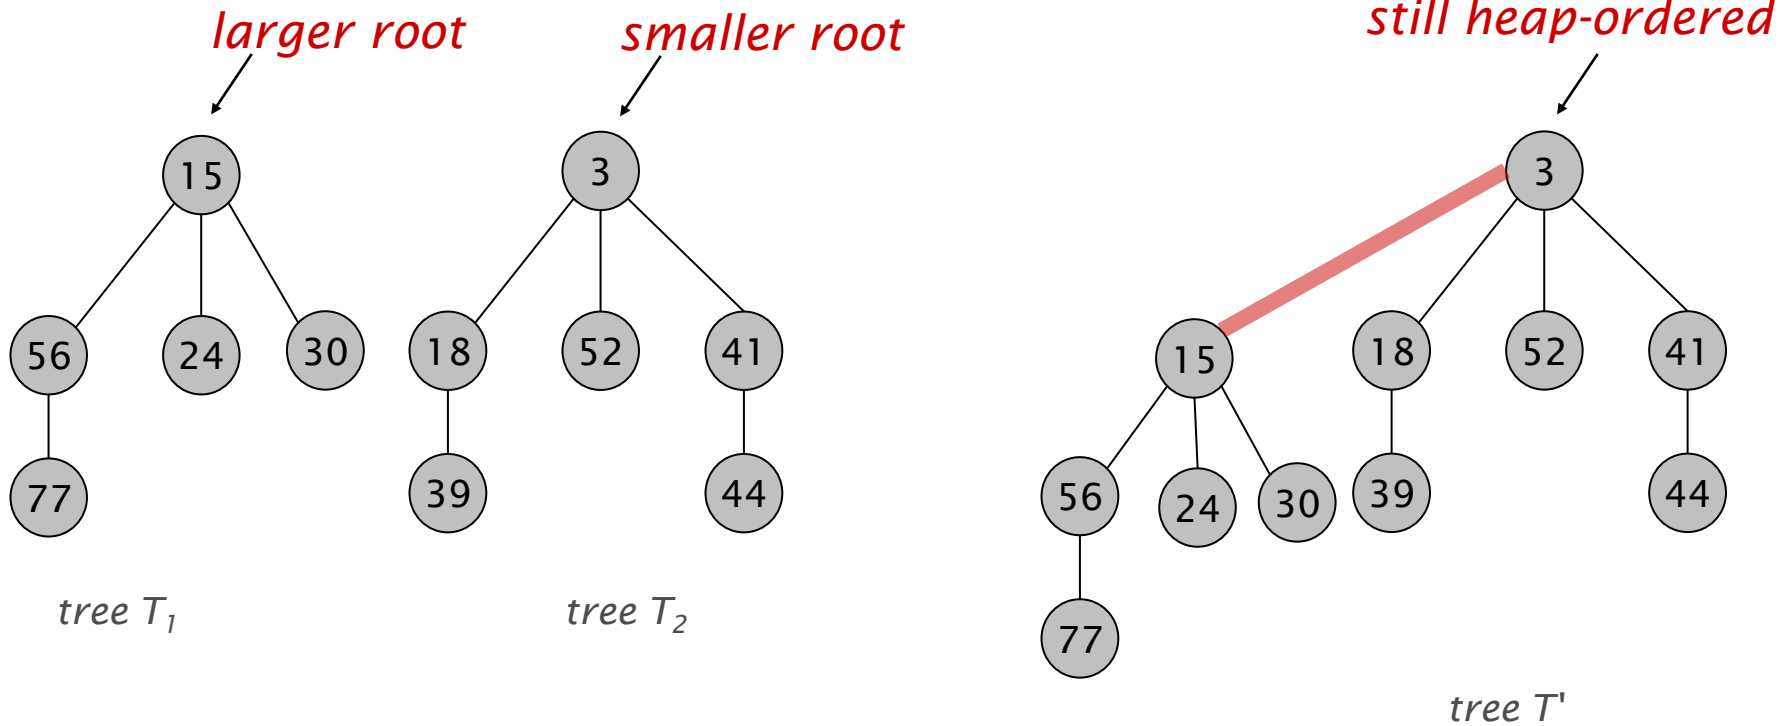
Only merge trees of the same rank

Recall $\text{rank}(T)$ is the number of children in the root of T

Make the larger root the child of the smaller root

Linking Operation


Linking operation. Make larger root be a child of smaller root.



Fibonacci Heaps: Delete Min

Delete min.


- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.


- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

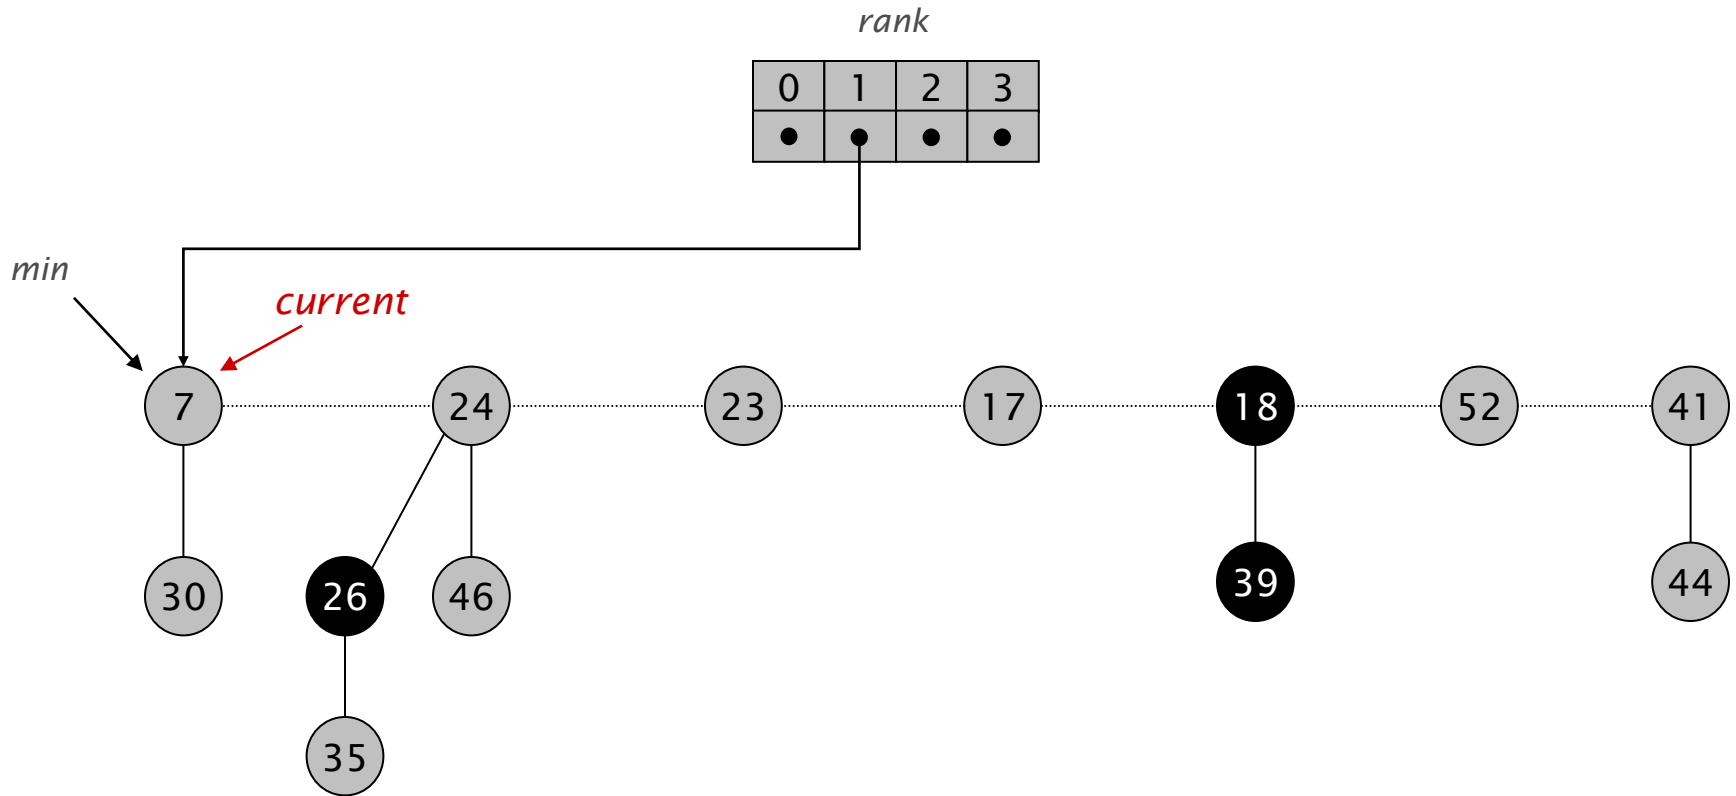
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

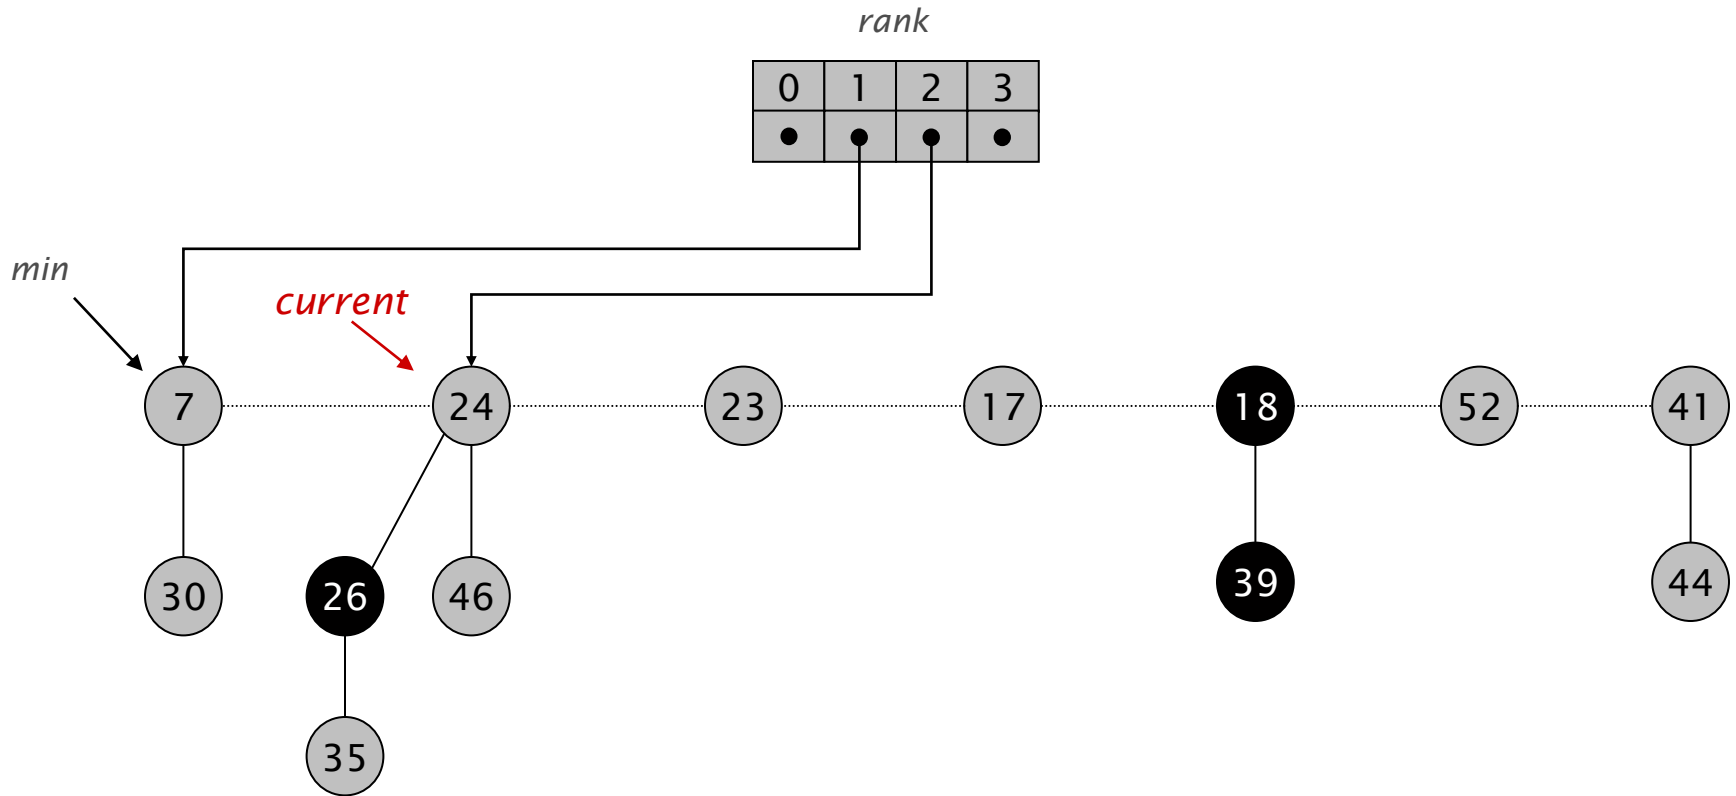
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

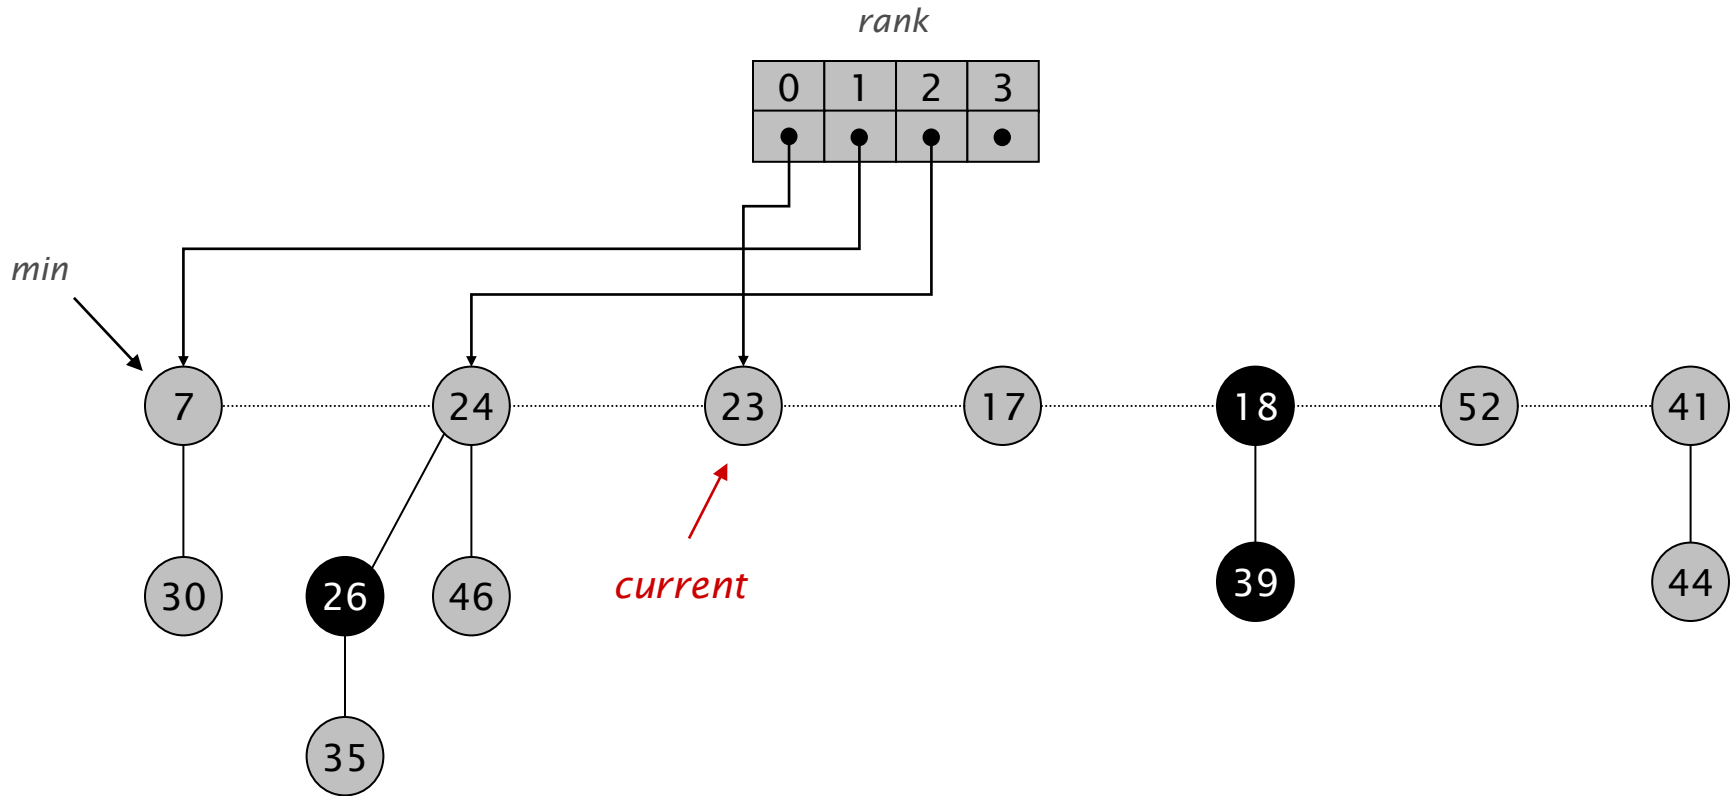
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

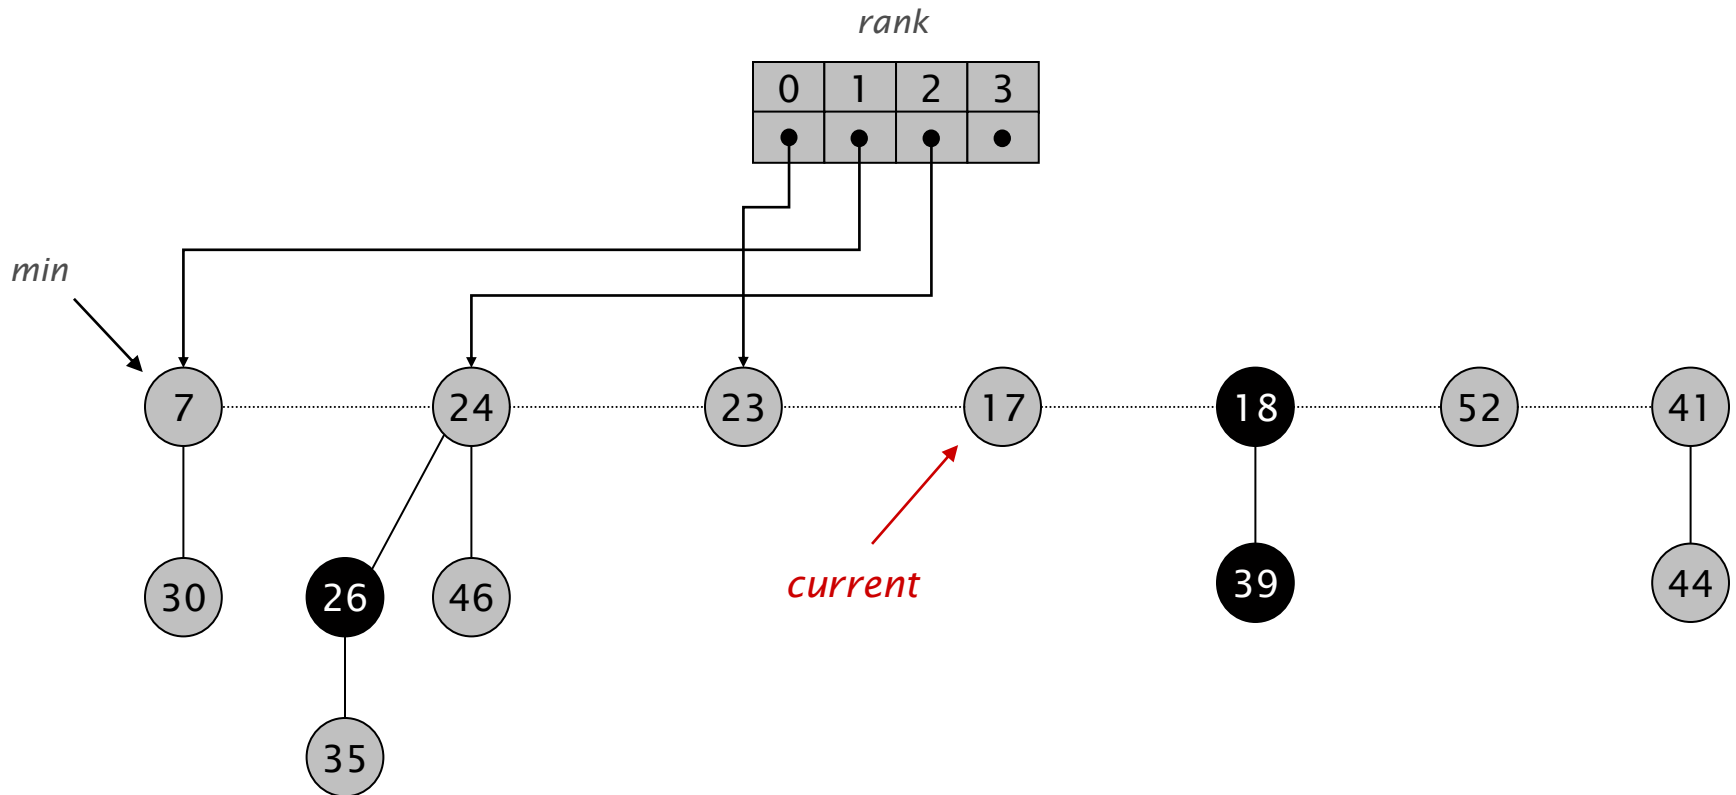
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.

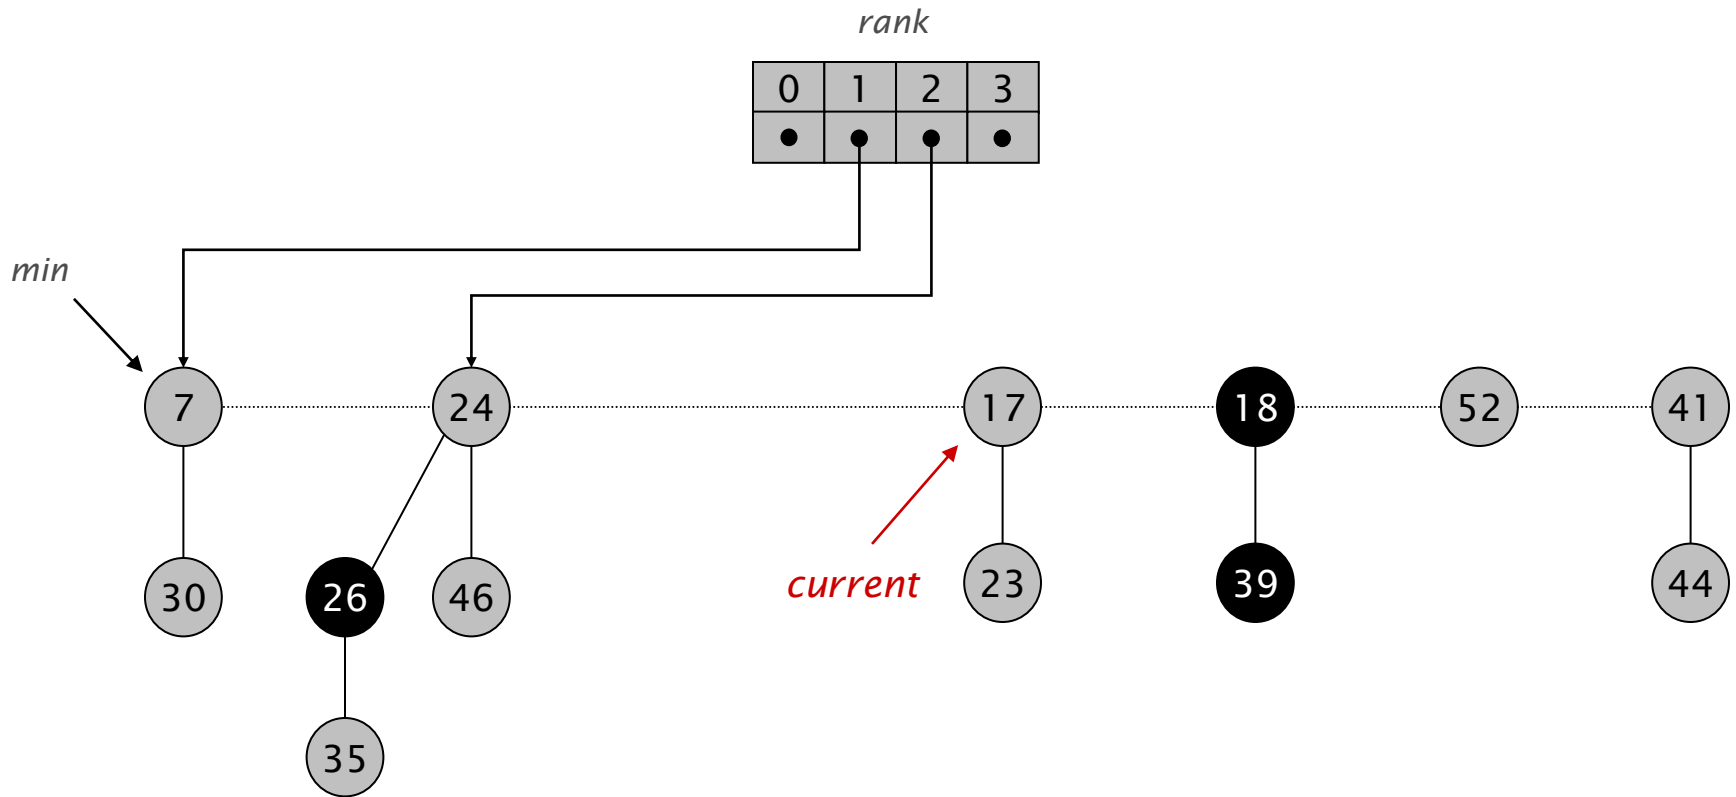


link 23 into 17

Fibonacci Heaps: Delete Min

Delete min.

- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.

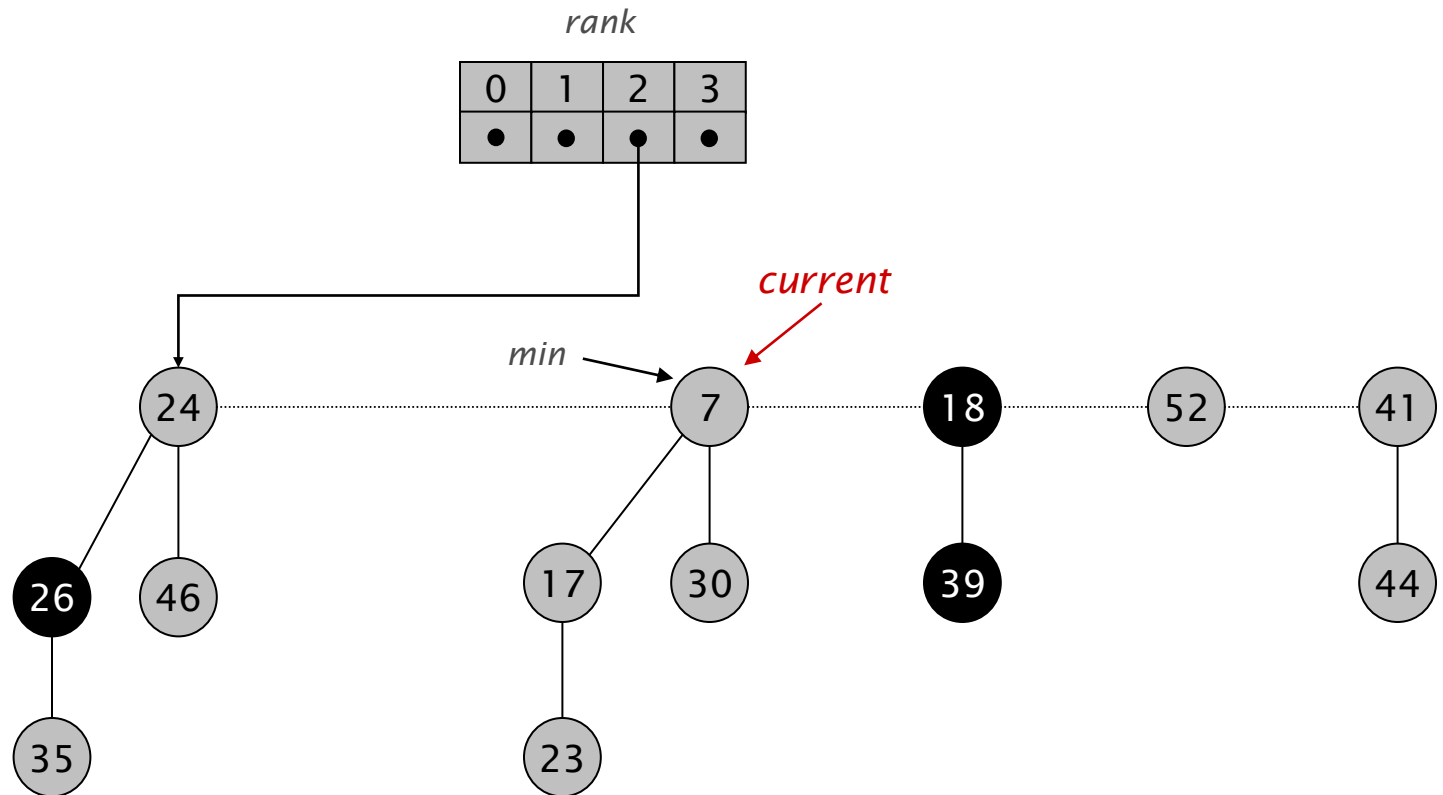


link 17 into 7

Fibonacci Heaps: Delete Min

Delete min.

- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.

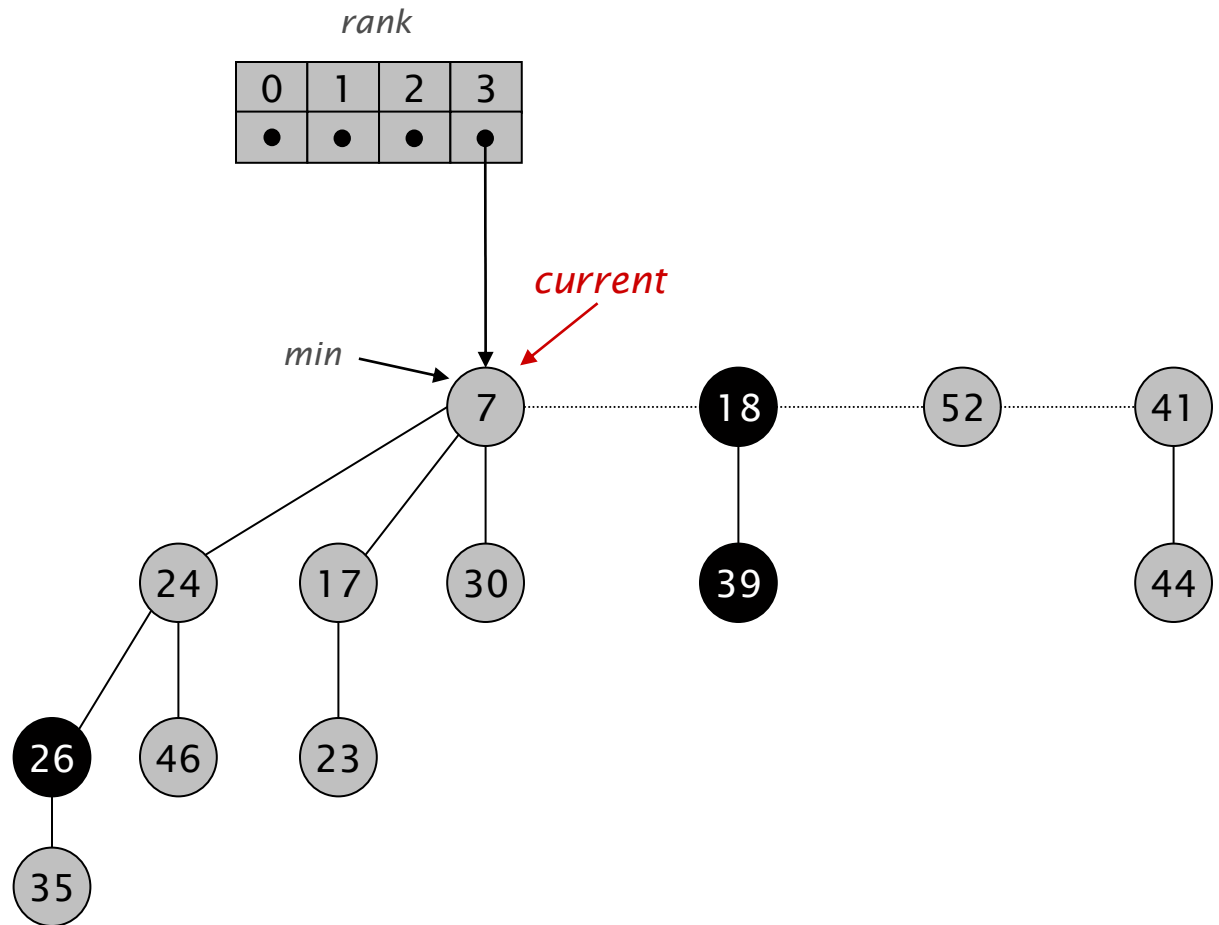


link 24 into 7

Fibonacci Heaps: Delete Min

Delete min.

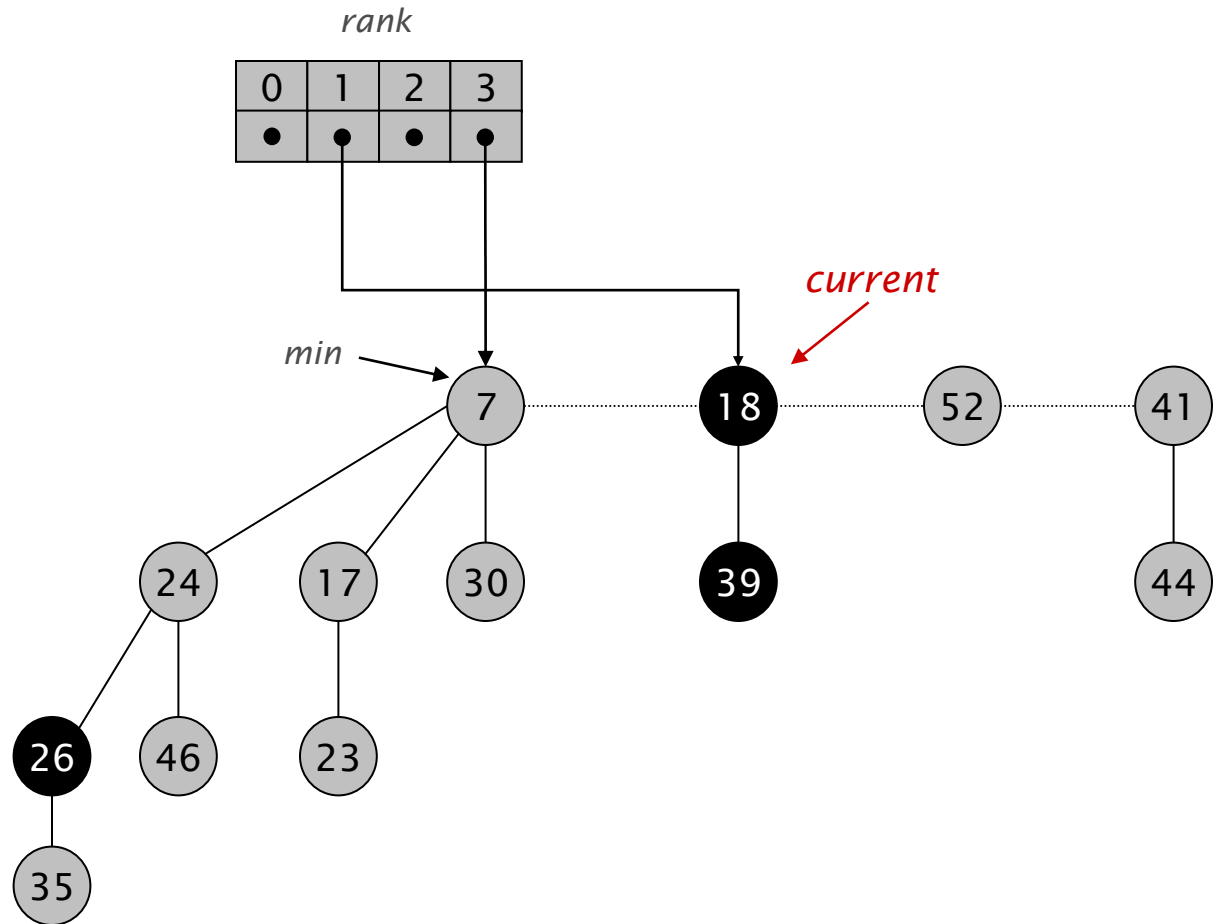
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

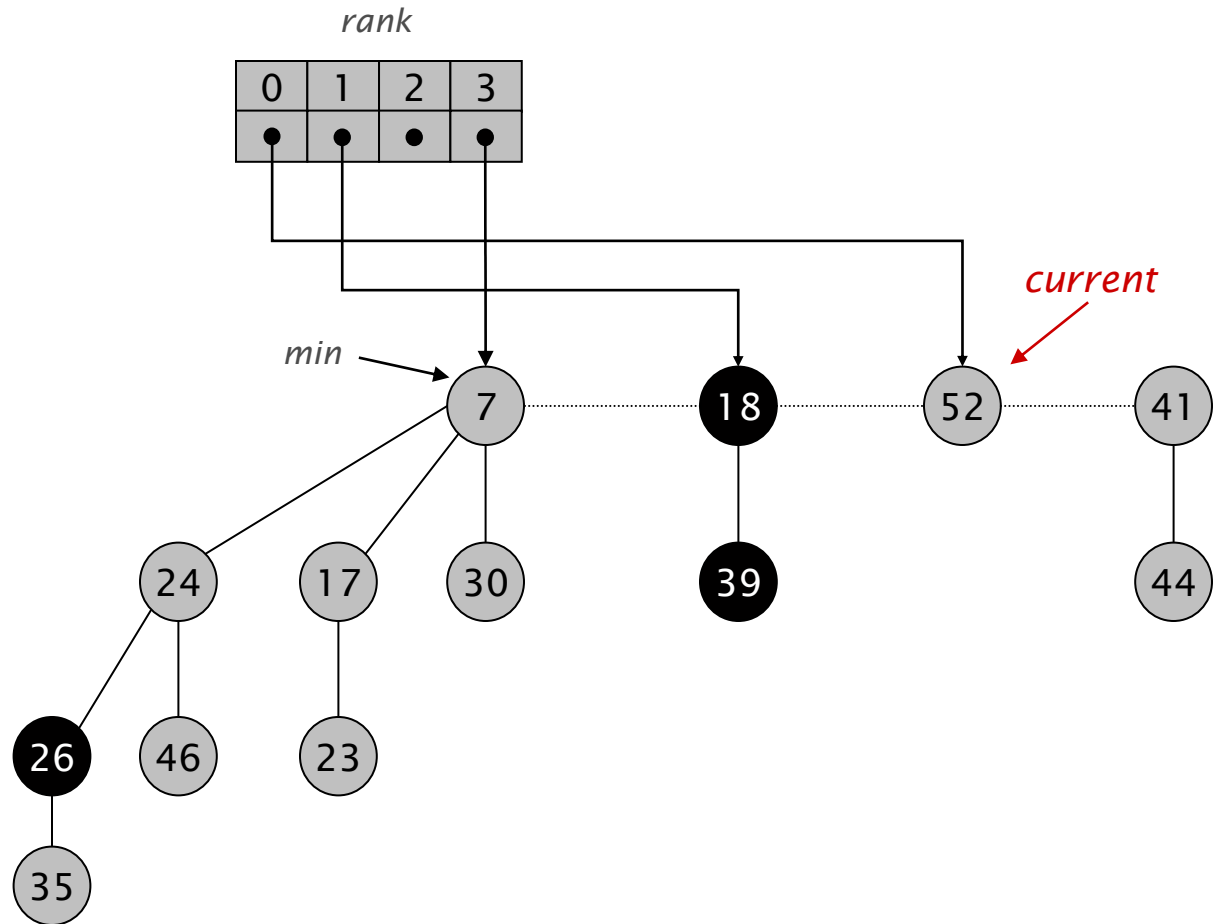
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.


- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

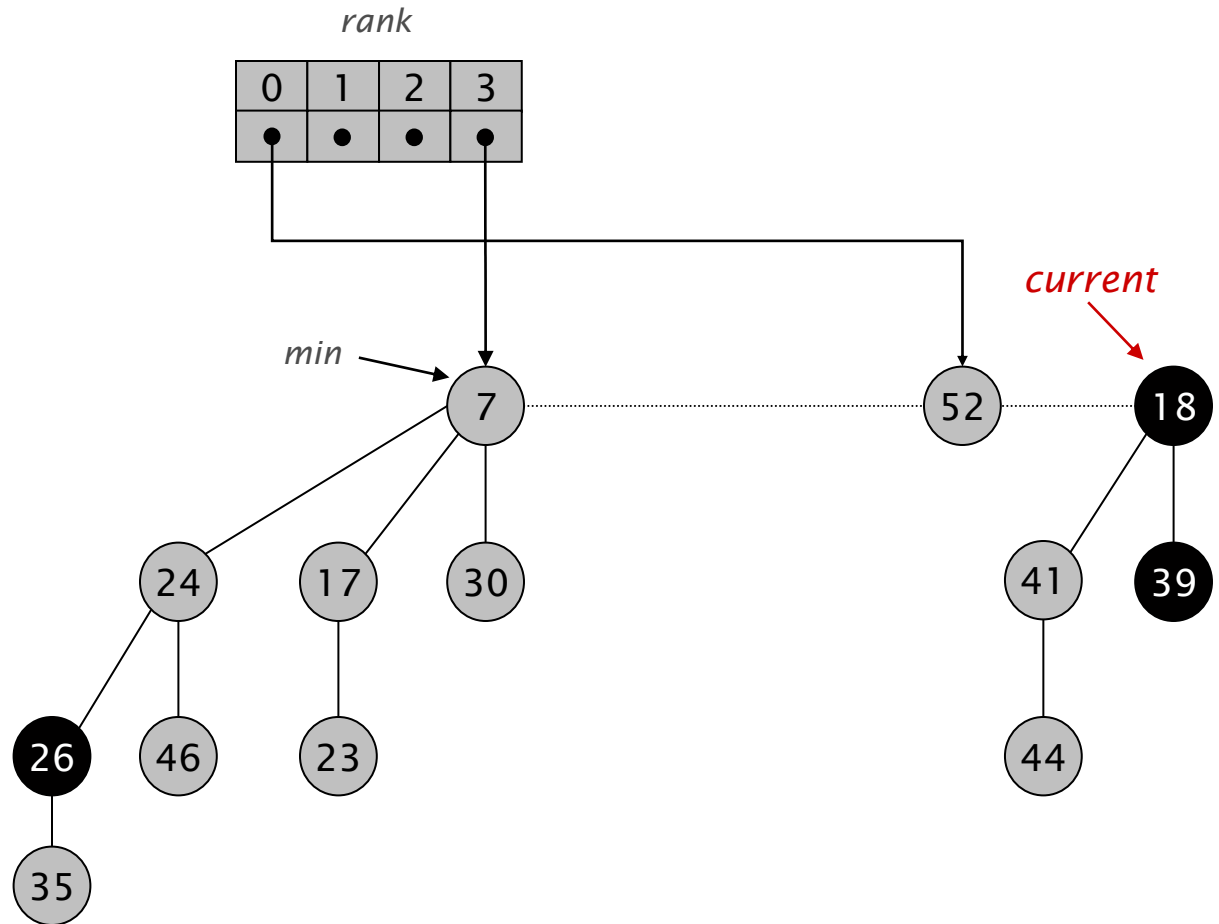
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

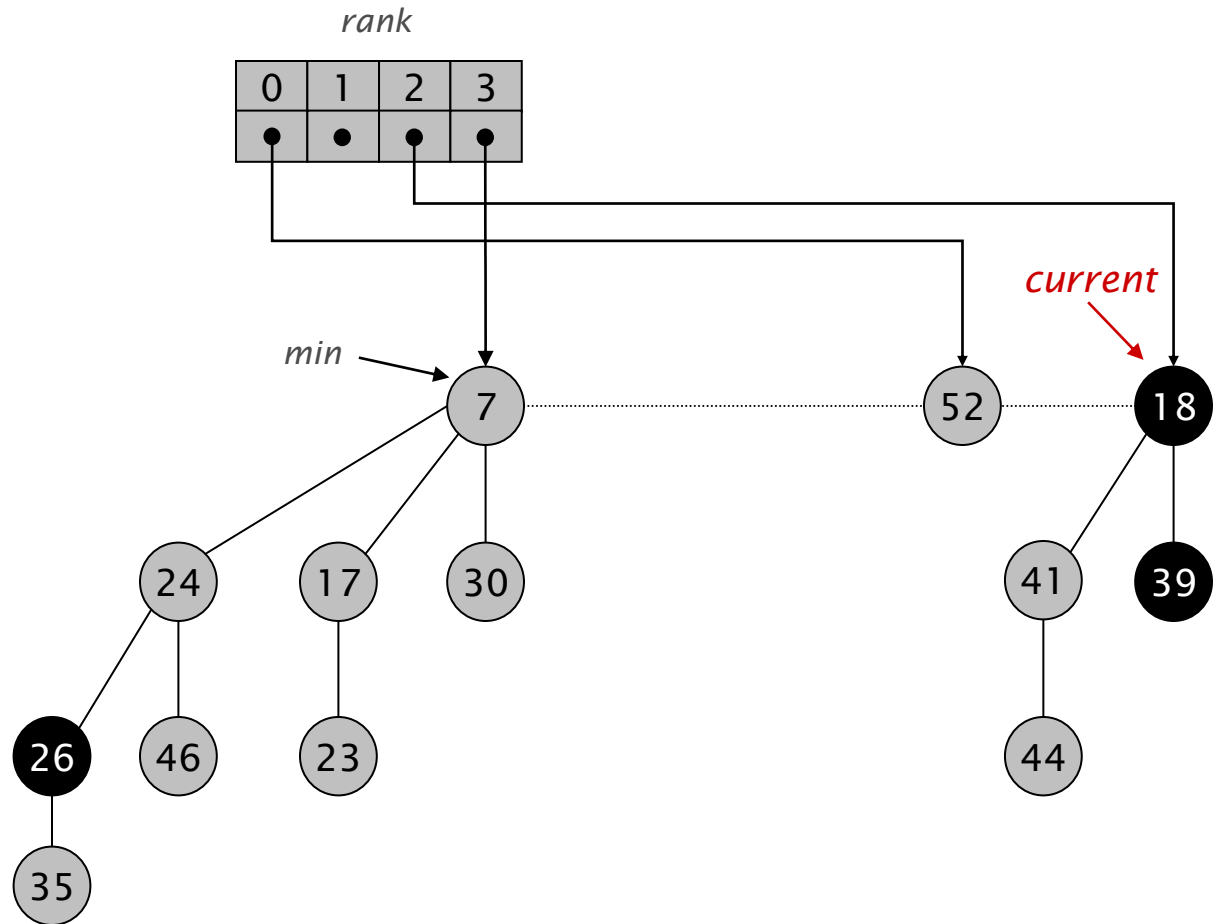
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

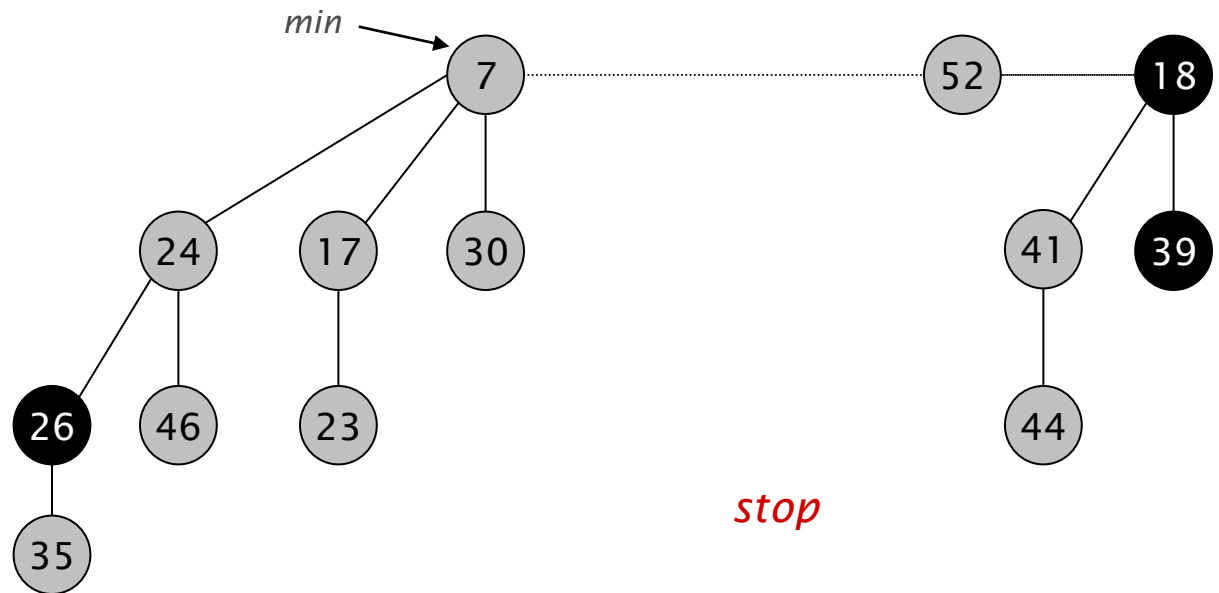
- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min

Delete min.

- Delete min; meld its children into root list; update min.
- Consolidate trees so that no two roots have same rank.



Fibonacci Heaps: Delete Min Analysis

Delete min.

$$\Phi(H) = \text{trees}(H) + 2 \cdot \text{marks}(H)$$

potential function

Actual cost. $O(\text{rank}(H)) + O(\text{trees}(H))$

- $O(\text{rank}(H))$ to meld min's children into root list.
- $O(\text{rank}(H)) + O(\text{trees}(H))$ to update min.
- $O(\text{rank}(H)) + O(\text{trees}(H))$ to consolidate trees.

Change in potential. $O(\text{rank}(H)) - \text{trees}(H)$

- *No change in marks(H)*
- $\text{trees}(H') \leq \text{rank}(H) + 1$ since no two trees have same rank.
- $\Delta\Phi(H) \leq \text{rank}(H) + 1 - \text{trees}(H)$.

Amortized cost. $O(\text{rank}(H))$


Fibonacci Heaps: Delete Min Analysis

Q. Is amortized cost of $O(\text{rank}(H))$ good?

A. Yes, if only *insert* and *delete-min* operations.

- In this case, all trees are binomial trees.
- This implies $\text{rank}(H) \leq \log n$.

we only link trees of equal rank



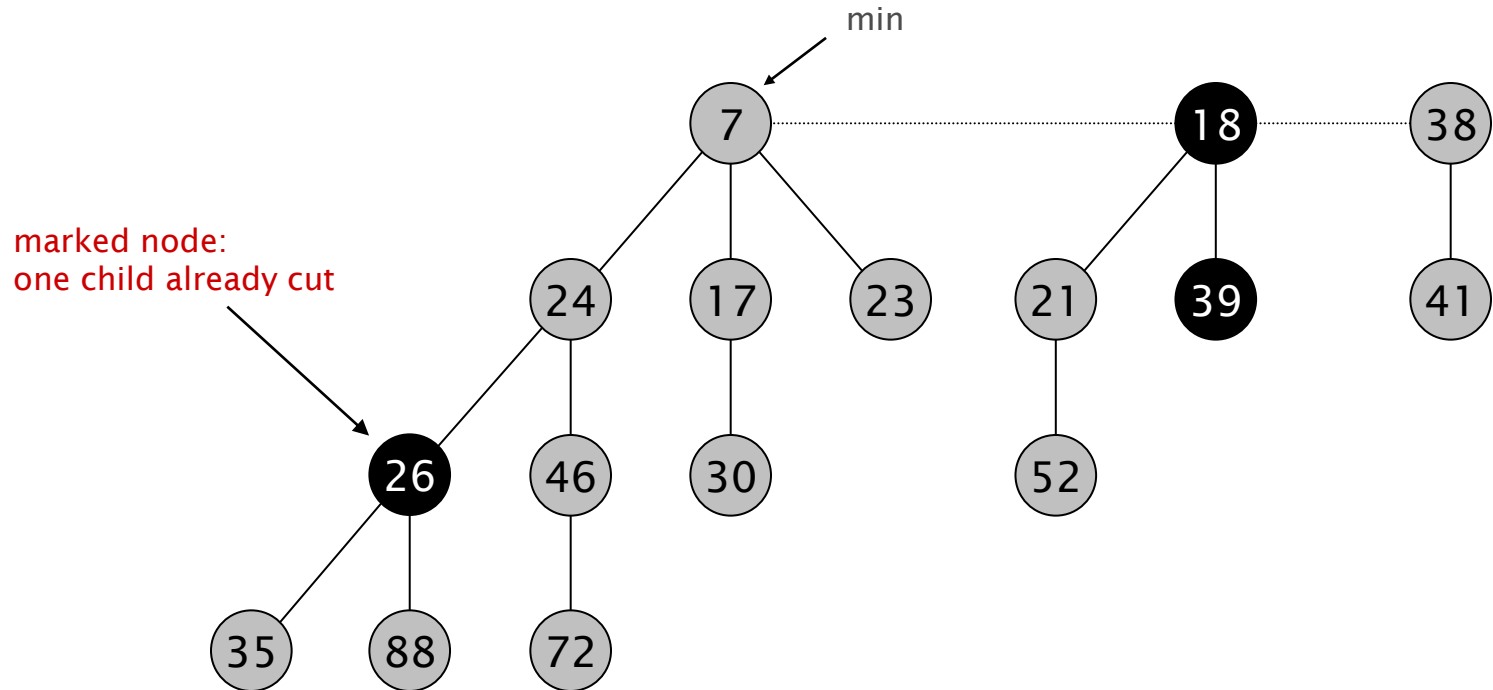
A. Yes, with Fib trees $\text{rank}(H) = O(\log n)$. Need to ensure that trees remain as Fib trees if not binomial trees

Decrease Key

Fibonacci Heaps: Decrease Key

Intuition for decreasing the key of node x .

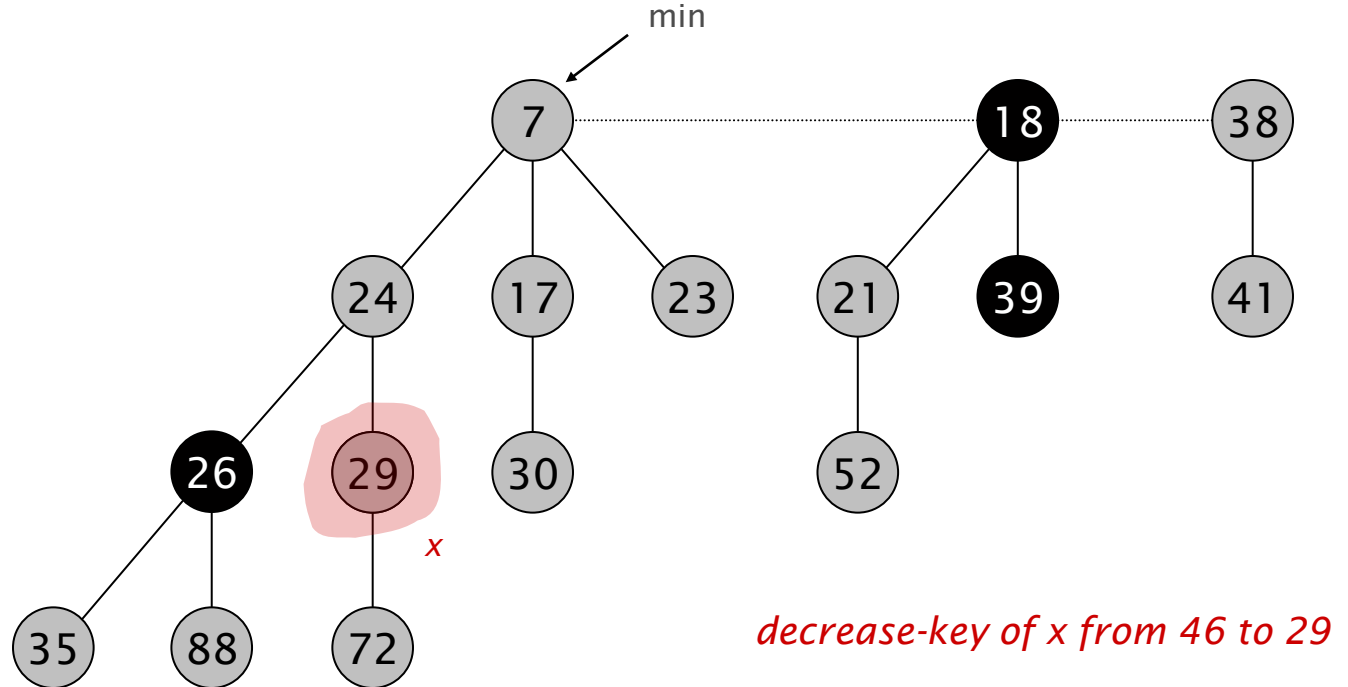
- If heap-order is not violated, just decrease the key of x .
- Otherwise, cut tree rooted at x and meld into root list.
- To keep trees flat: as soon as a node has its second child cut, cut it off and meld into root list (and unmark it).



Fibonacci Heaps: Decrease Key

Case 1. [heap order not violated]

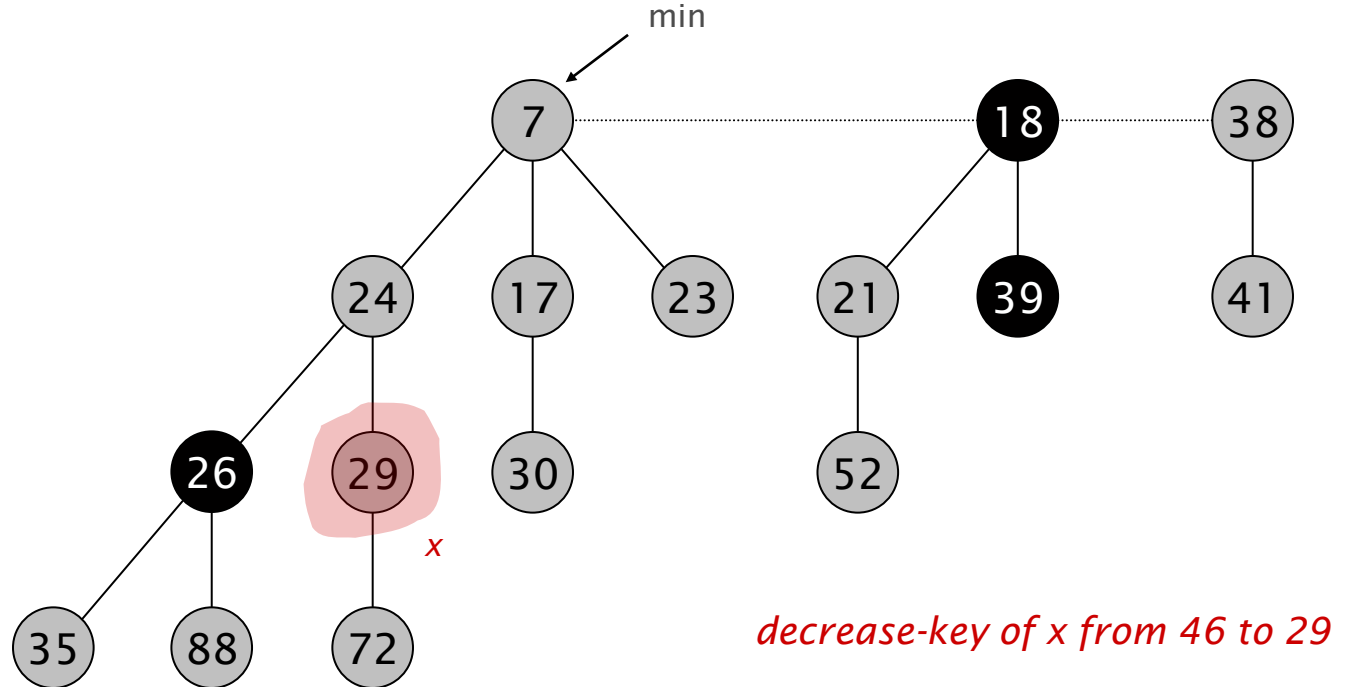
- Decrease key of x .
- Change heap min pointer (if necessary).



Fibonacci Heaps: Decrease Key

Case 1. [heap order not violated]


- Decrease key of x .
- Change heap min pointer (if necessary).



Fibonacci Heaps: Decrease Key

Case 2a. [heap order violated]


- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2a. [heap order violated]

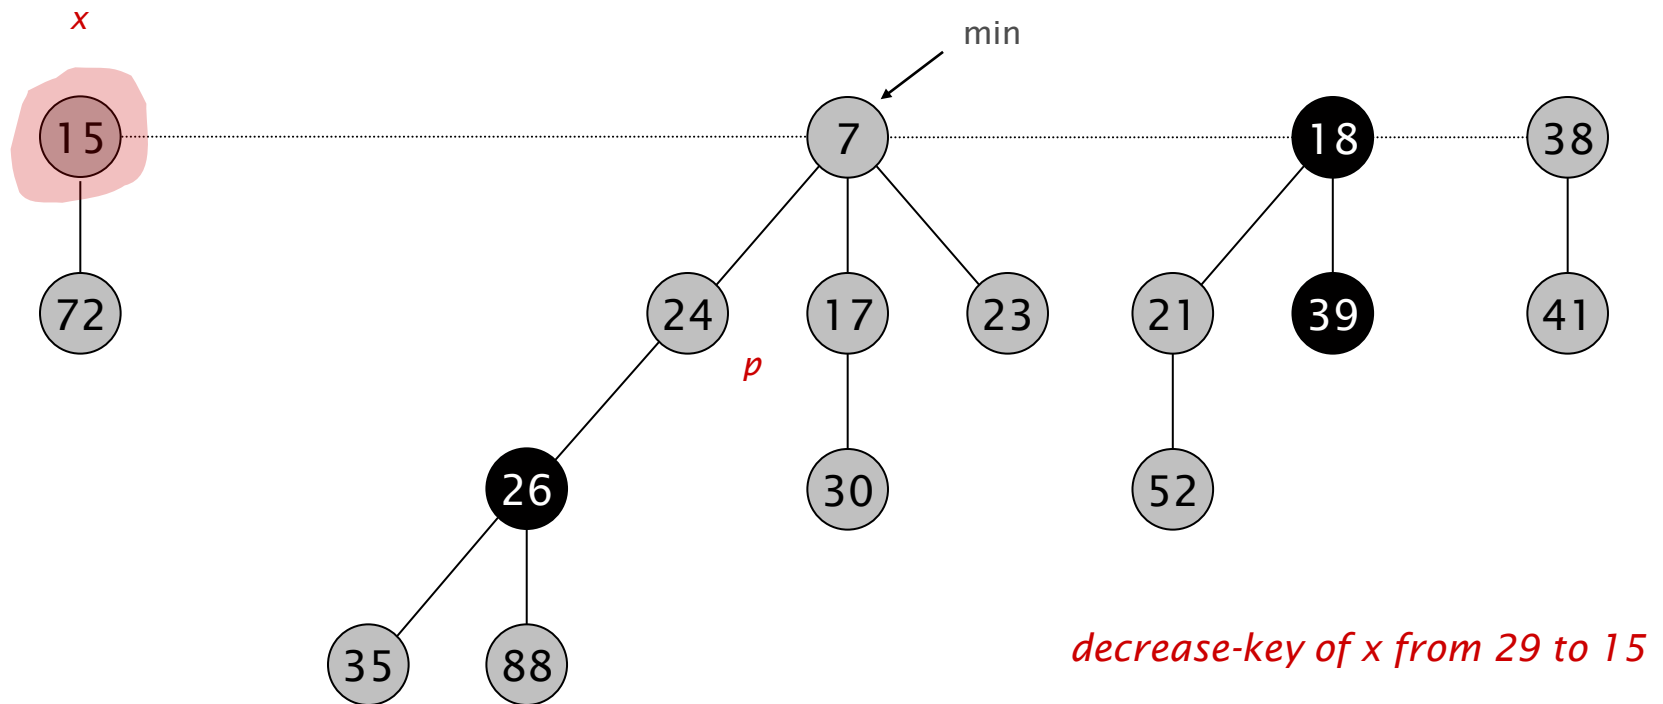
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2a. [heap order violated]

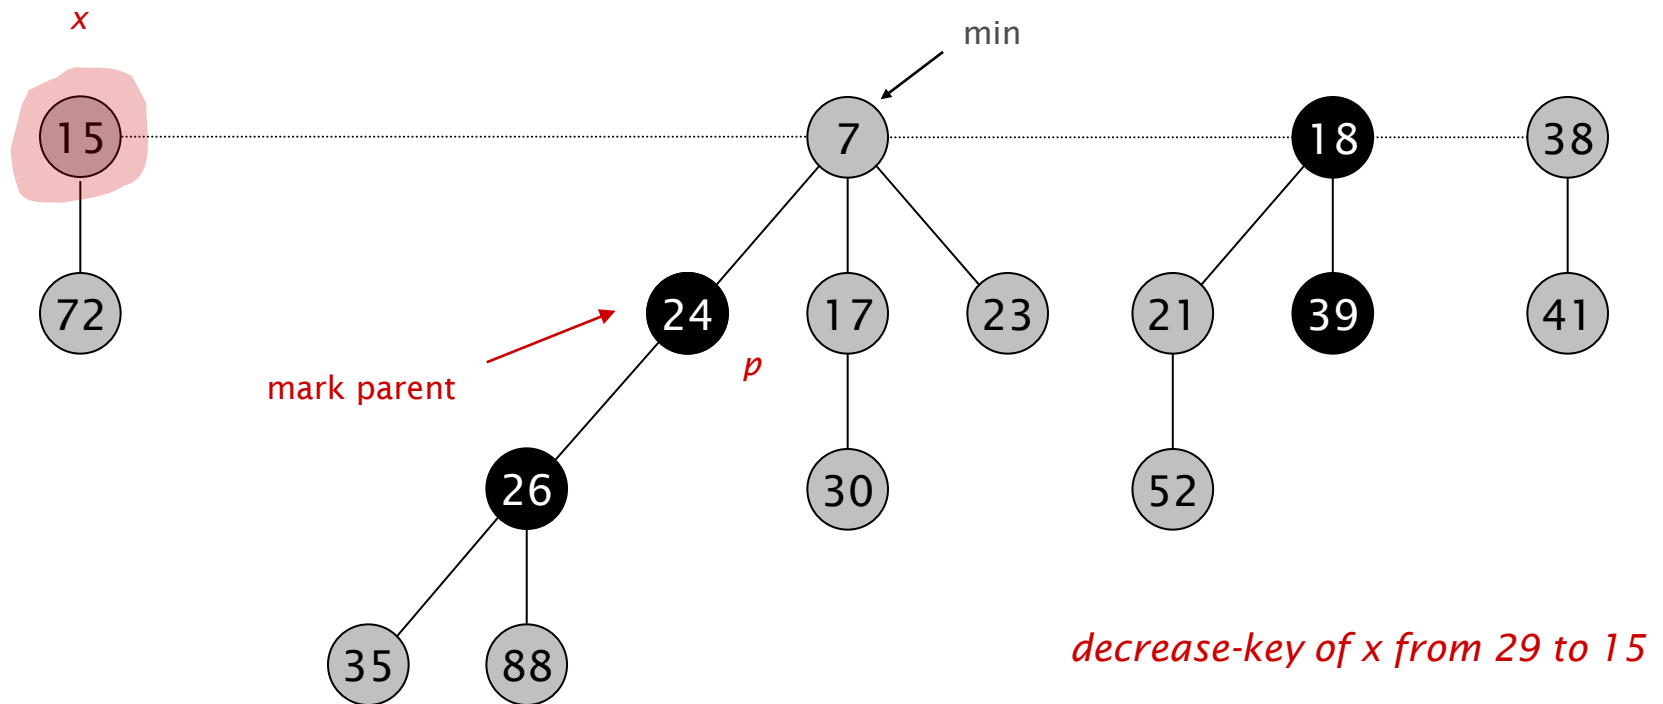
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2a. [heap order violated]

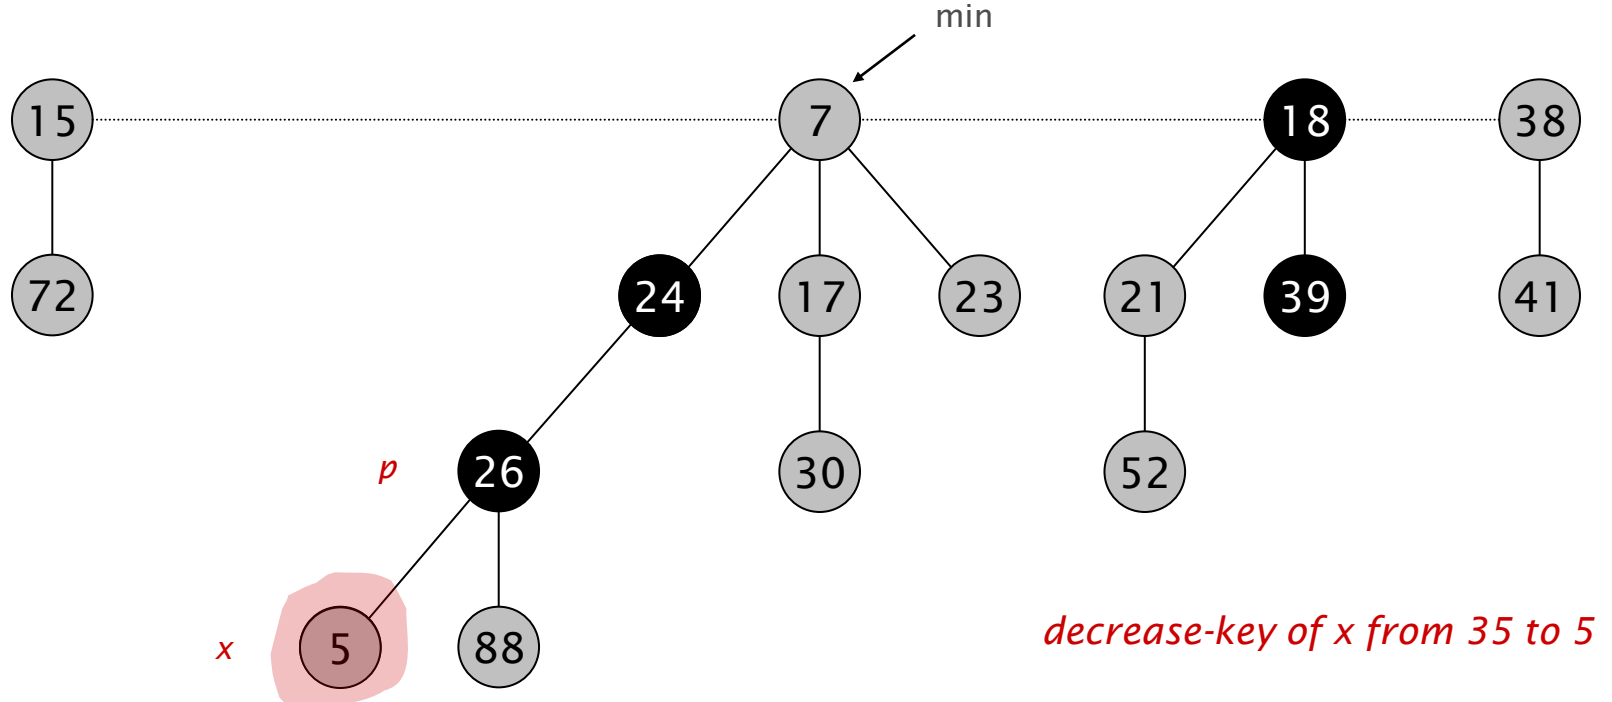
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

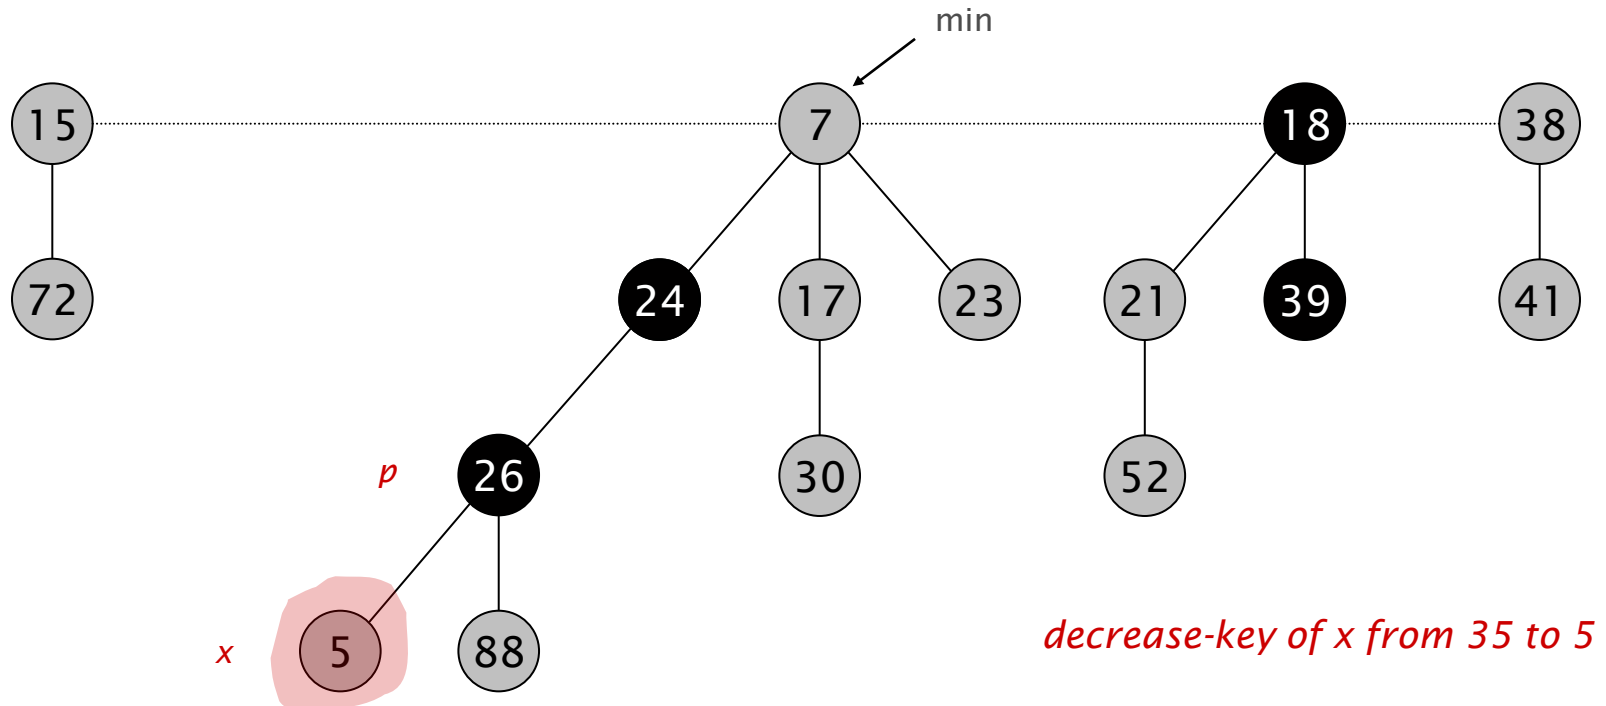
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

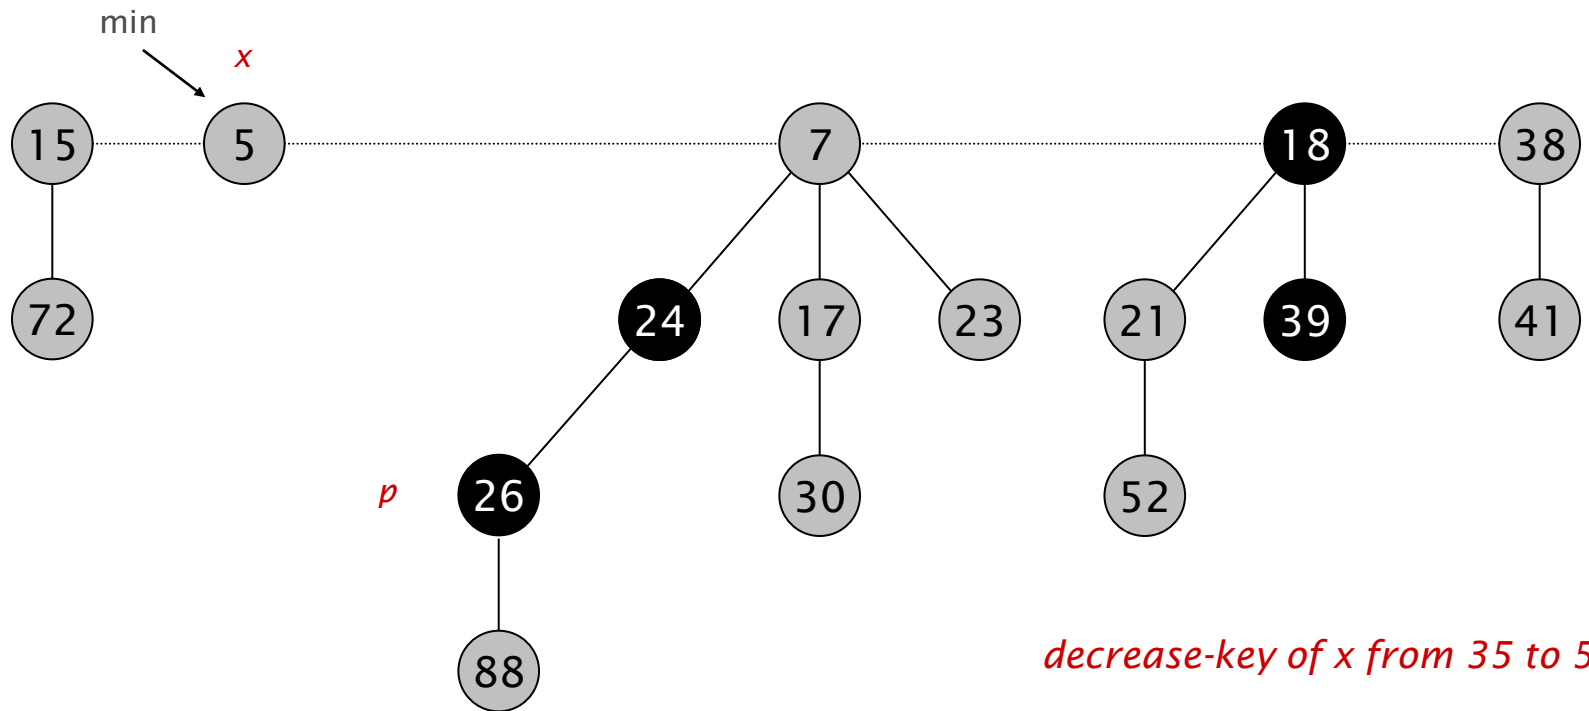
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

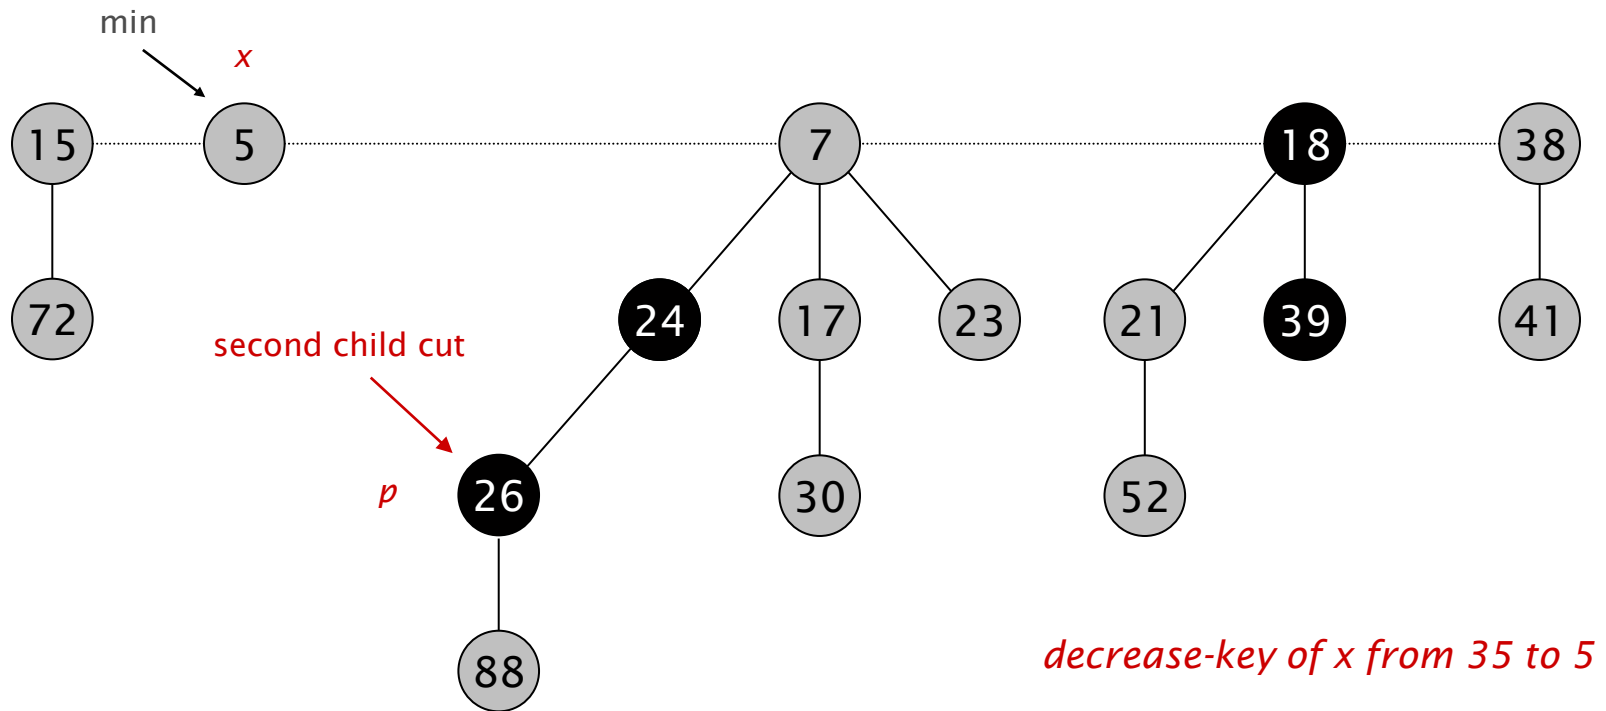
- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]


- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).

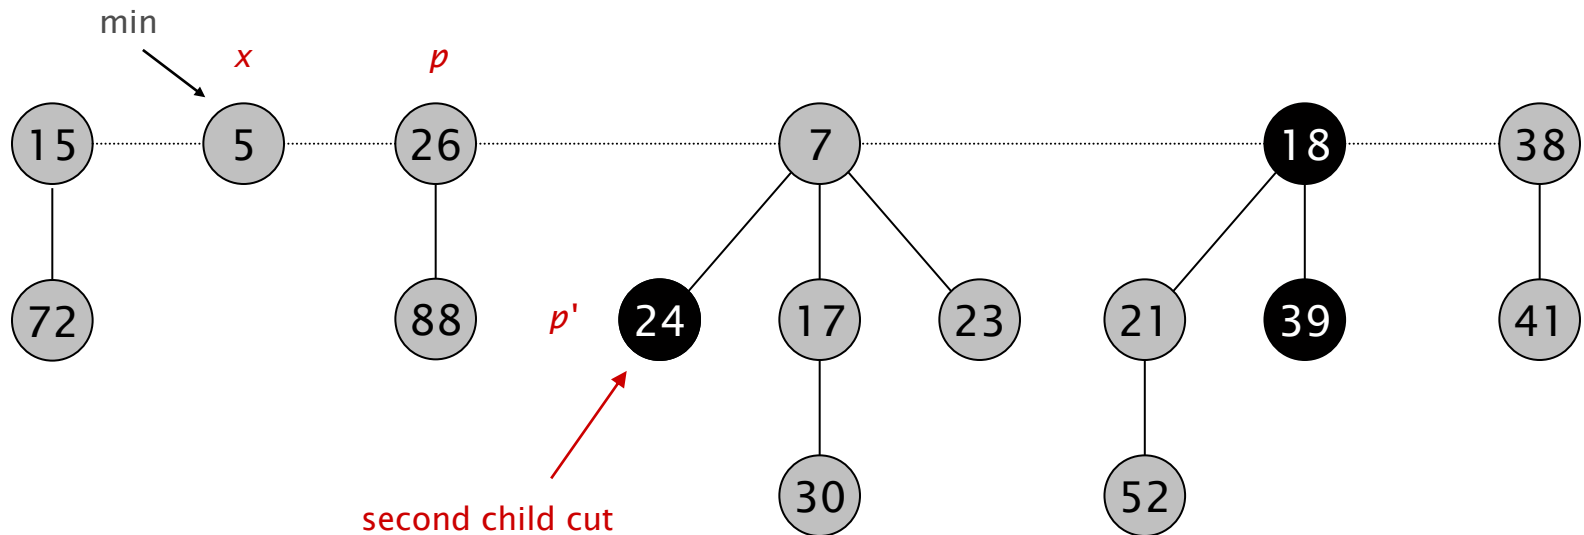


decrease-key of x from 35 to 5

Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).

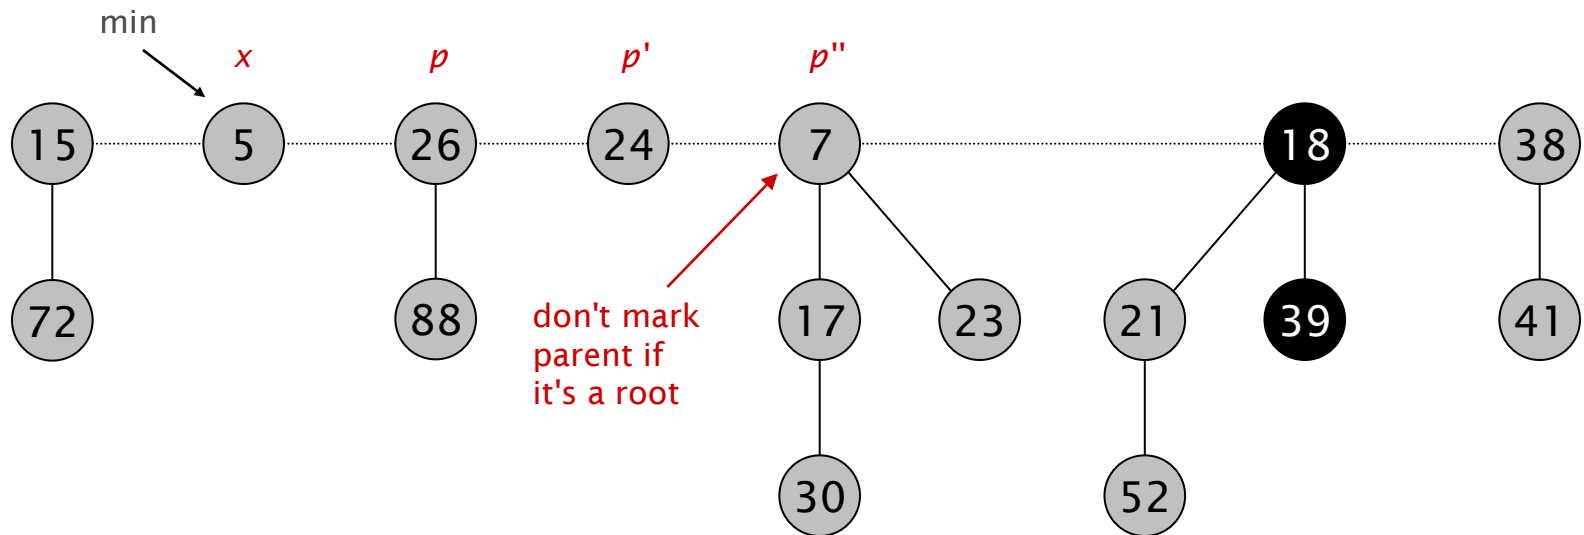


decrease-key of x from 35 to 5

Fibonacci Heaps: Decrease Key

Case 2b. [heap order violated]

- Decrease key of x .
- Cut tree rooted at x , meld into root list, and unmark.
- If parent p of x is unmarked (hasn't yet lost a child), mark it; Otherwise, cut p , meld into root list, and unmark (and do so recursively for all ancestors that lose a second child).



decrease-key of x from 35 to 5

Fibonacci Heaps: Decrease Key Analysis

Decrease-key.

$$\Phi(H) = \text{trees}(H) + 2 \cdot \text{marks}(H)$$

potential function

Actual cost. $O(c)$, where c is the number of cuts needed

- $O(1)$ time for changing the key.
- $O(1)$ time for each of c cuts, plus melding into root list.

Change in potential. $O(1) - c$

- $\text{trees}(H') = \text{trees}(H) + c.$
- $\text{marks}(H') \leq \text{marks}(H) - c + 2.$
- $\Delta\Phi \leq c + 2 \cdot (-c + 2) = 4 - c.$

Amortized cost. $O(1)$

Analysis

- Need to show that all the trees are Fib trees
- Need to prove the two properties of Fib trees

Proving Fib Trees

Lemma. Let x be a node with rank k , and let y_1, \dots, y_k denote the children of x in the order in which they were linked to x . Then:

$$\text{rank}(y_i) \geq \begin{cases} 0, & \text{if } i = 1 \\ i - 2, & \text{if } i \geq 2 \end{cases}$$

Proof.

- When y_i is linked to x , y_1, \dots, y_{i-1} already linked to x ,
 \Rightarrow In that step, $\text{rank}(x) \geq i - 1$
 $\Rightarrow \text{rank}(y_i) \geq i - 1$ since we only link nodes of equal degree
- Since then, y_i has lost at most one child
 - otherwise it would have been cut from x
- Thus, $\text{rank}(y_i) \geq i - 2$

Properties of Fib Trees

Define: $\text{size}(T)$ as the number of nodes in T

$D(T) = \max$ degree of the nodes in T

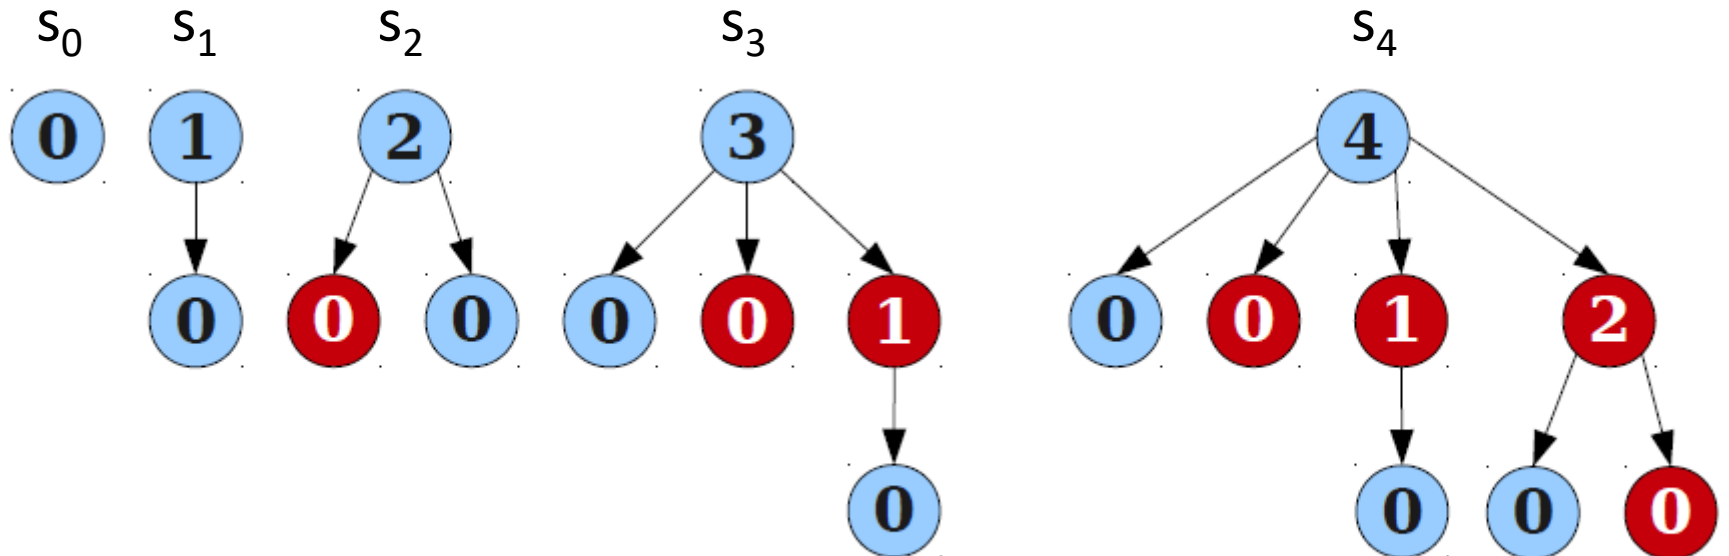
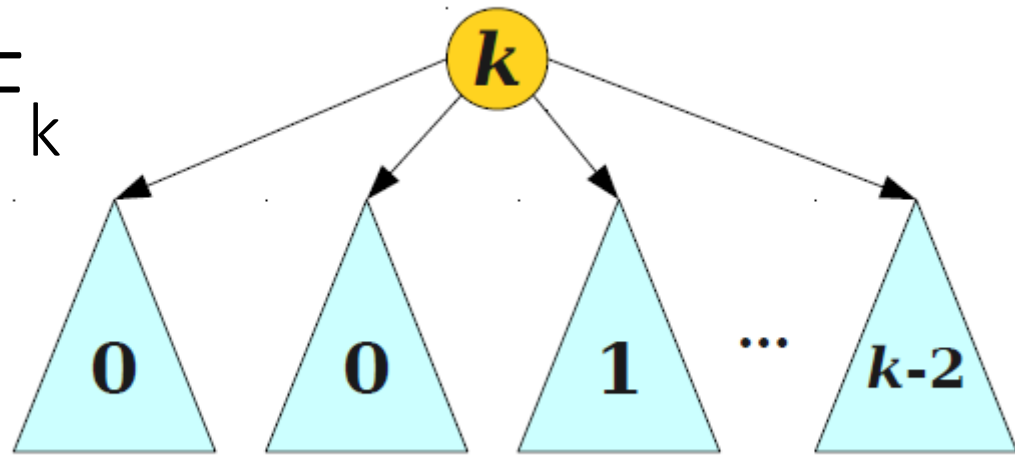
Lemma: For the Fib tree F_k

1. $\text{size}(F_k) \geq \phi^k$ where $\phi = (1 + \sqrt{5}) / 2$
2. $D(F_k) \leq \log_\phi(\text{size}(F_k))$

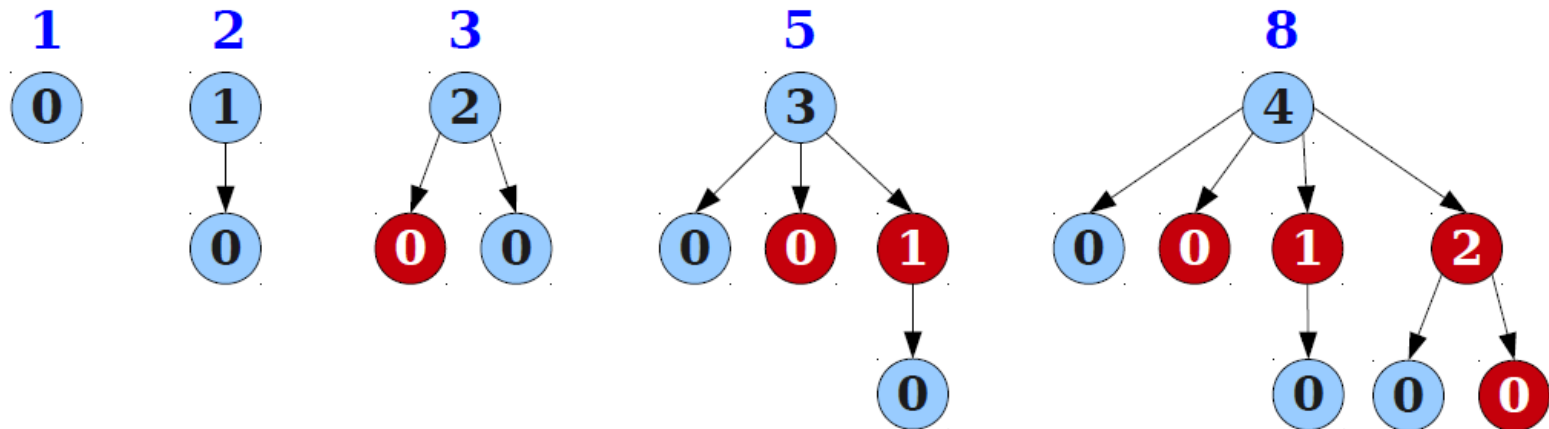
Homework exercise

Bounding Size of F_k

Let s_k be the smallest possible size of F_k



Bounding Size of F_k

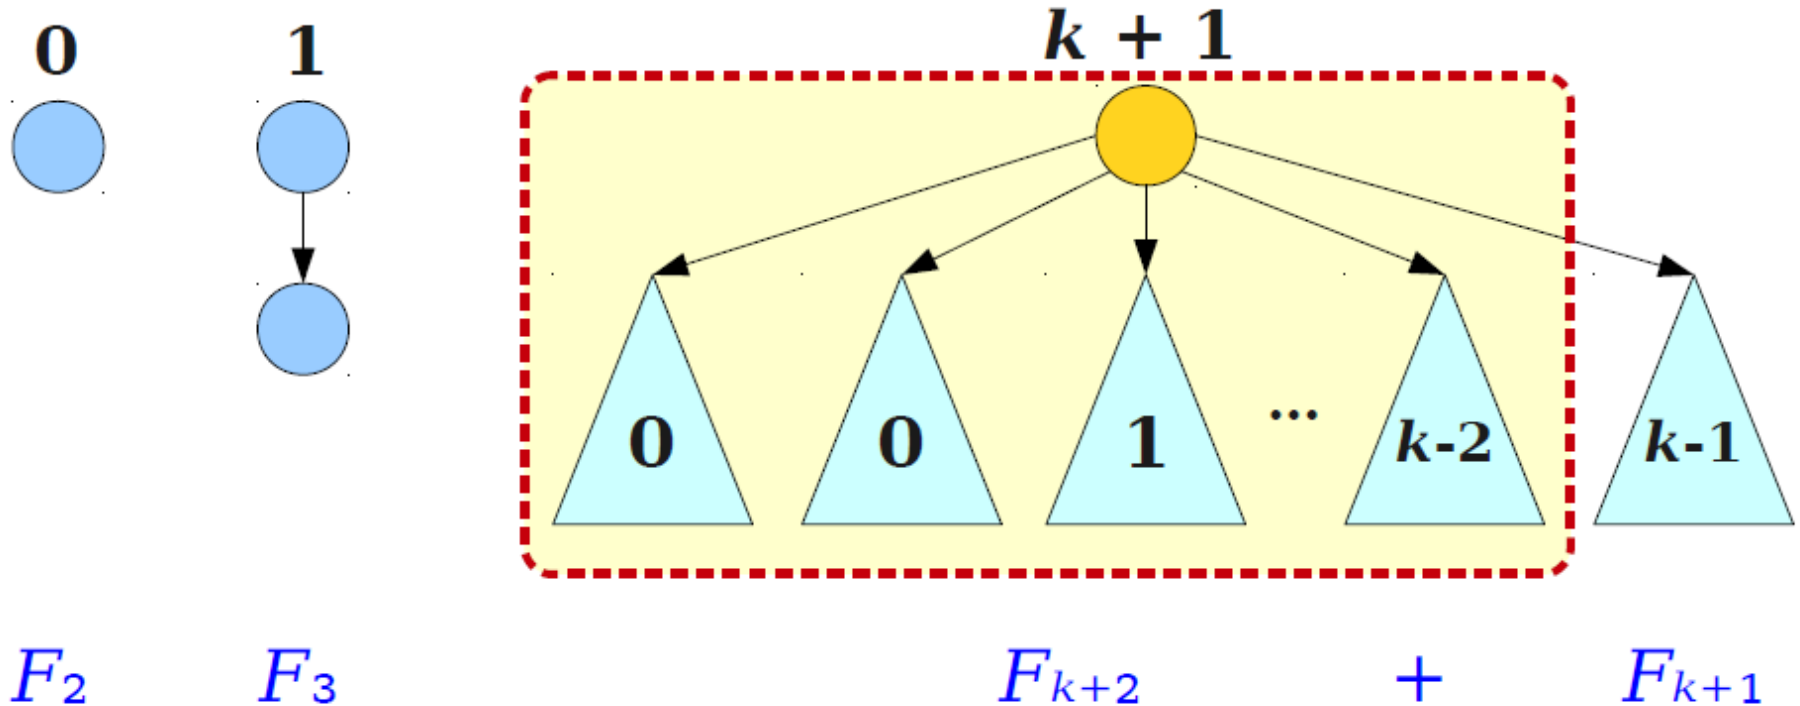


Claim: The minimum number of nodes in a tree of rank k is F_{k+2}

Bounding Size of F_k

Theorem: The number of nodes in a Fib tree of rank k is F_{k+2} .

Proof: Induction.



Bounding the Rank of Nodes in F_k


- **Fact:** For $n \geq 2$, we have $F_n \geq \phi^{n-2}$, where ϕ is the golden ratio: $\phi \approx 1.61803398875\dots$
- **Claim:** In our modified data structure, we have $\text{rank}(T) = O(\log n)$.
- **Proof:** In a tree of rank k , there are at least $F_{k+2} \geq \phi^k$ nodes. Therefore, the maximum rank of a tree in our data structure is $\log_\phi n = O(\log n)$.

Fibonacci Heaps: Summary

Operation	Linked List	Binary Heap	Binomial Heap	Fibonacci Heap [†]	Relaxed Heap
<i>make-heap</i>	1	1	1	1	1
<i>is-empty</i>	1	1	1	1	1
<i>insert</i>	1	$\log n$	$\log n$	1	1
<i>delete-min</i>	n	$\log n$	$\log n$	$\log n$	$\log n$
<i>decrease-key</i>	n	$\log n$	$\log n$	1	1
<i>delete</i>	n	$\log n$	$\log n$	$\log n$	$\log n$
<i>union</i>	1	n	$\log n$	1	1
<i>find-min</i>	n	1	$\log n$	1	1

n = number of elements in priority queue

Runtime of
Dijkstra's/Prim's
Algorithm



Thank You